



TensorFlow

Machine Learning and Deep Learning with TensorFlow

History

1st Generation: DistBelief



Dean. et al. 2011

Major Products:

- Inception (Image Categorization)
- Google Search
- Google Translate
- Google Photos

2nd Generation: TensorFlow

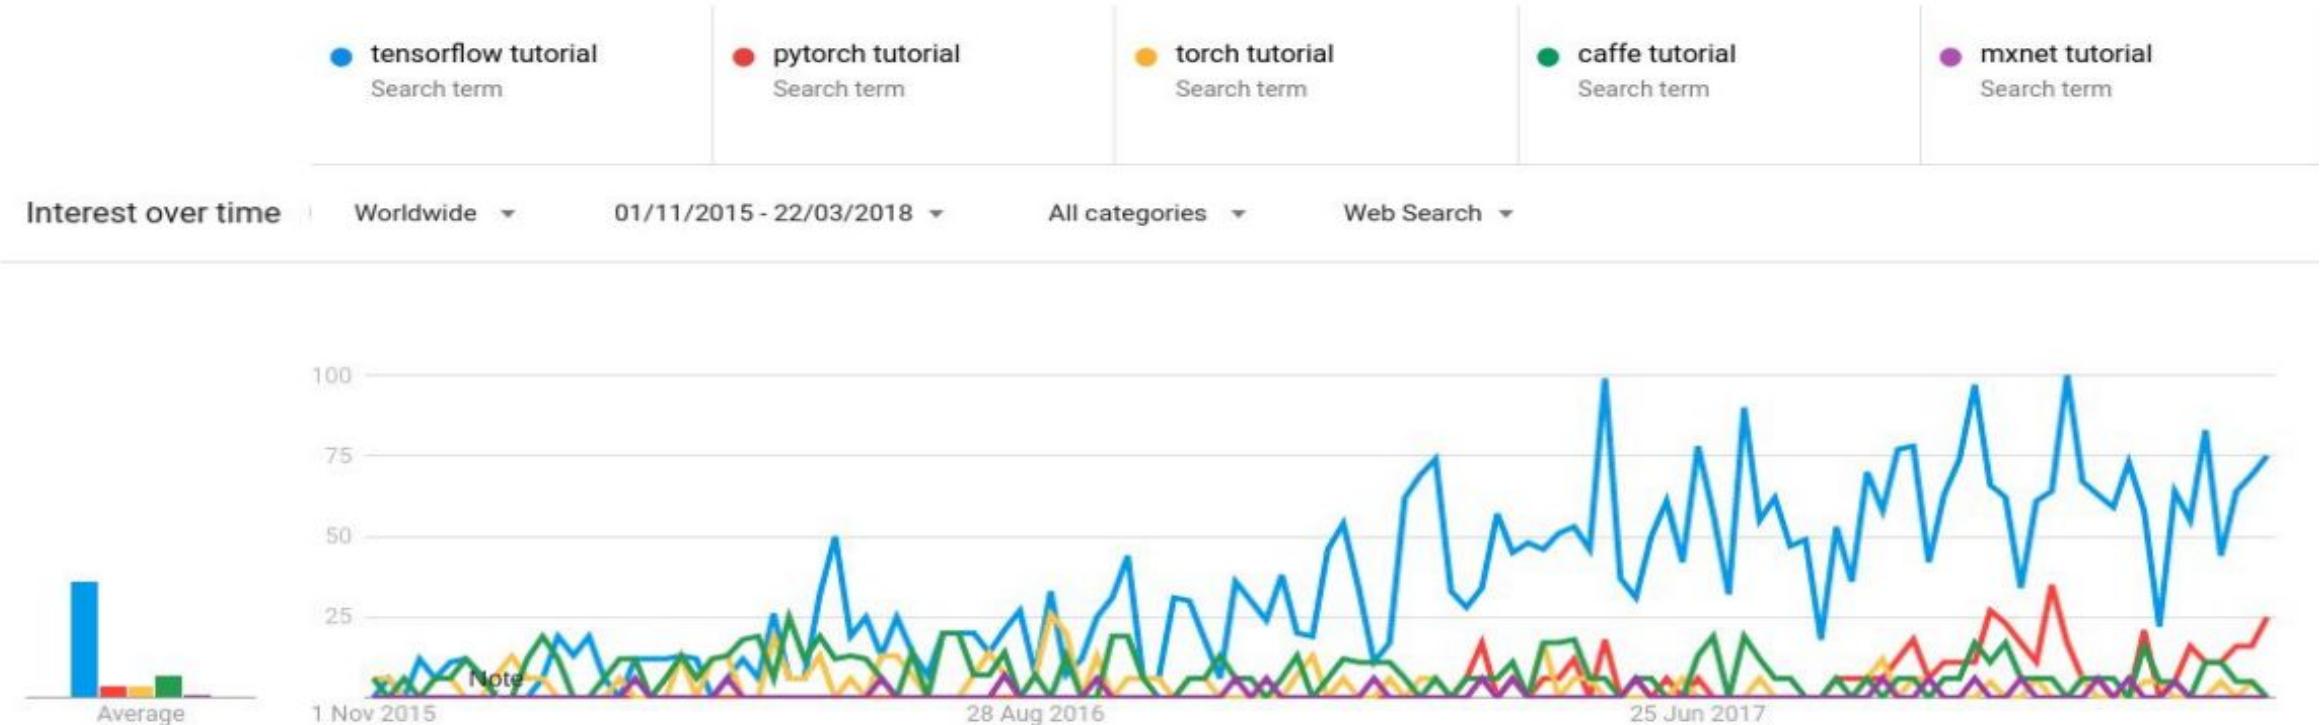


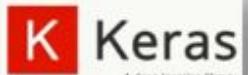
- Dean. et al. 2015 (1st November)
- Most of DistBelief users at Google have already switched to TensorFlow

Jeff Dean – One of the main developer of DistBelief and TensorFlow

Why TensorFlow

- “TensorFlow™ is an open source software library for numerical computation using data flow graphs.”
- One of many frameworks for deep learning computations
- Scalable and flexible





A deep learning library

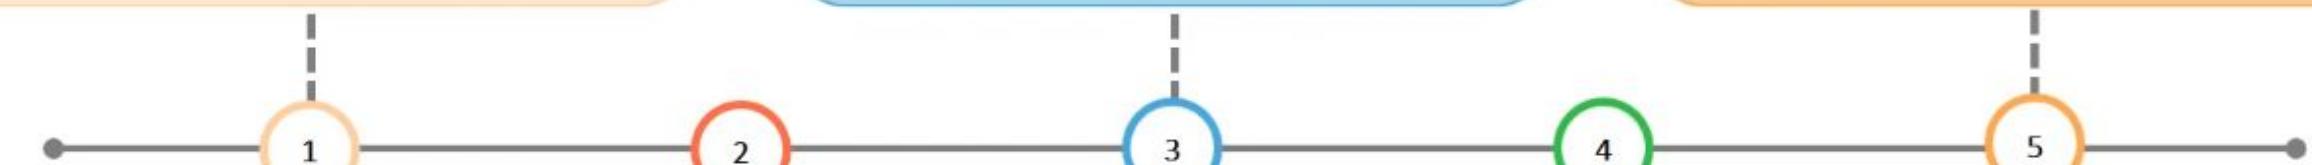
- Developed by Francois Chollet
- Open source library written in Python

theano

- Developed by University of Montreal
- Written in Python



- Developed by Google Brain Team
- Written in C++, Python and CUDA



- Developed by Skymind engineering team and DeepLearning4J community
- Written in C++ and Java



- Created by Ronan Collobert, Koray kavukcuoglu, Clement Farabet
- Written in Python

Companies using TensorFlow

- Google
- OpenAI
- DeepMind
- Snapchat
- Uber
- Airbus
- eBay
- Dropbox
- A bunch of startups...

What are Tensors?

- Tensors are the standard way of representing data in Tensor Flow (deep learning).
- Tensors are multidimensional arrays, an extension of two-dimensional tables(matrices) to data with higher dimension

't'
'e'
'n'
's'
'o'
'r'

*Tensor of
dimension[1]*

3	1	4	1
5	9	2	6
5	3	5	8
9	7	9	3
2	3	8	4
6	2	6	4

*Tensor of
dimensions[2]*

2	1	8	8	1
2	8	4	5	0
2	3	9	0	4
2	5	6	2	8

*Tensor of
dimensions[3]*

Tensors Rank

Rank	Math Entity	Python Example
0	Scalar(magnitude only)	S=434
1	Vector(magnitude and direction)	V=[1.1,2.2,3.3]
2	Matrix(table of numbers)	M=[[1,2,3],[4,5,6],[7,8,9]]
3	3-Tensor(cube of the numbers)	T=[[[2],[4],[6],[8],[10],[12]], [[14],[16],[18]]]
n	N-Tensor(you get the idea)	...

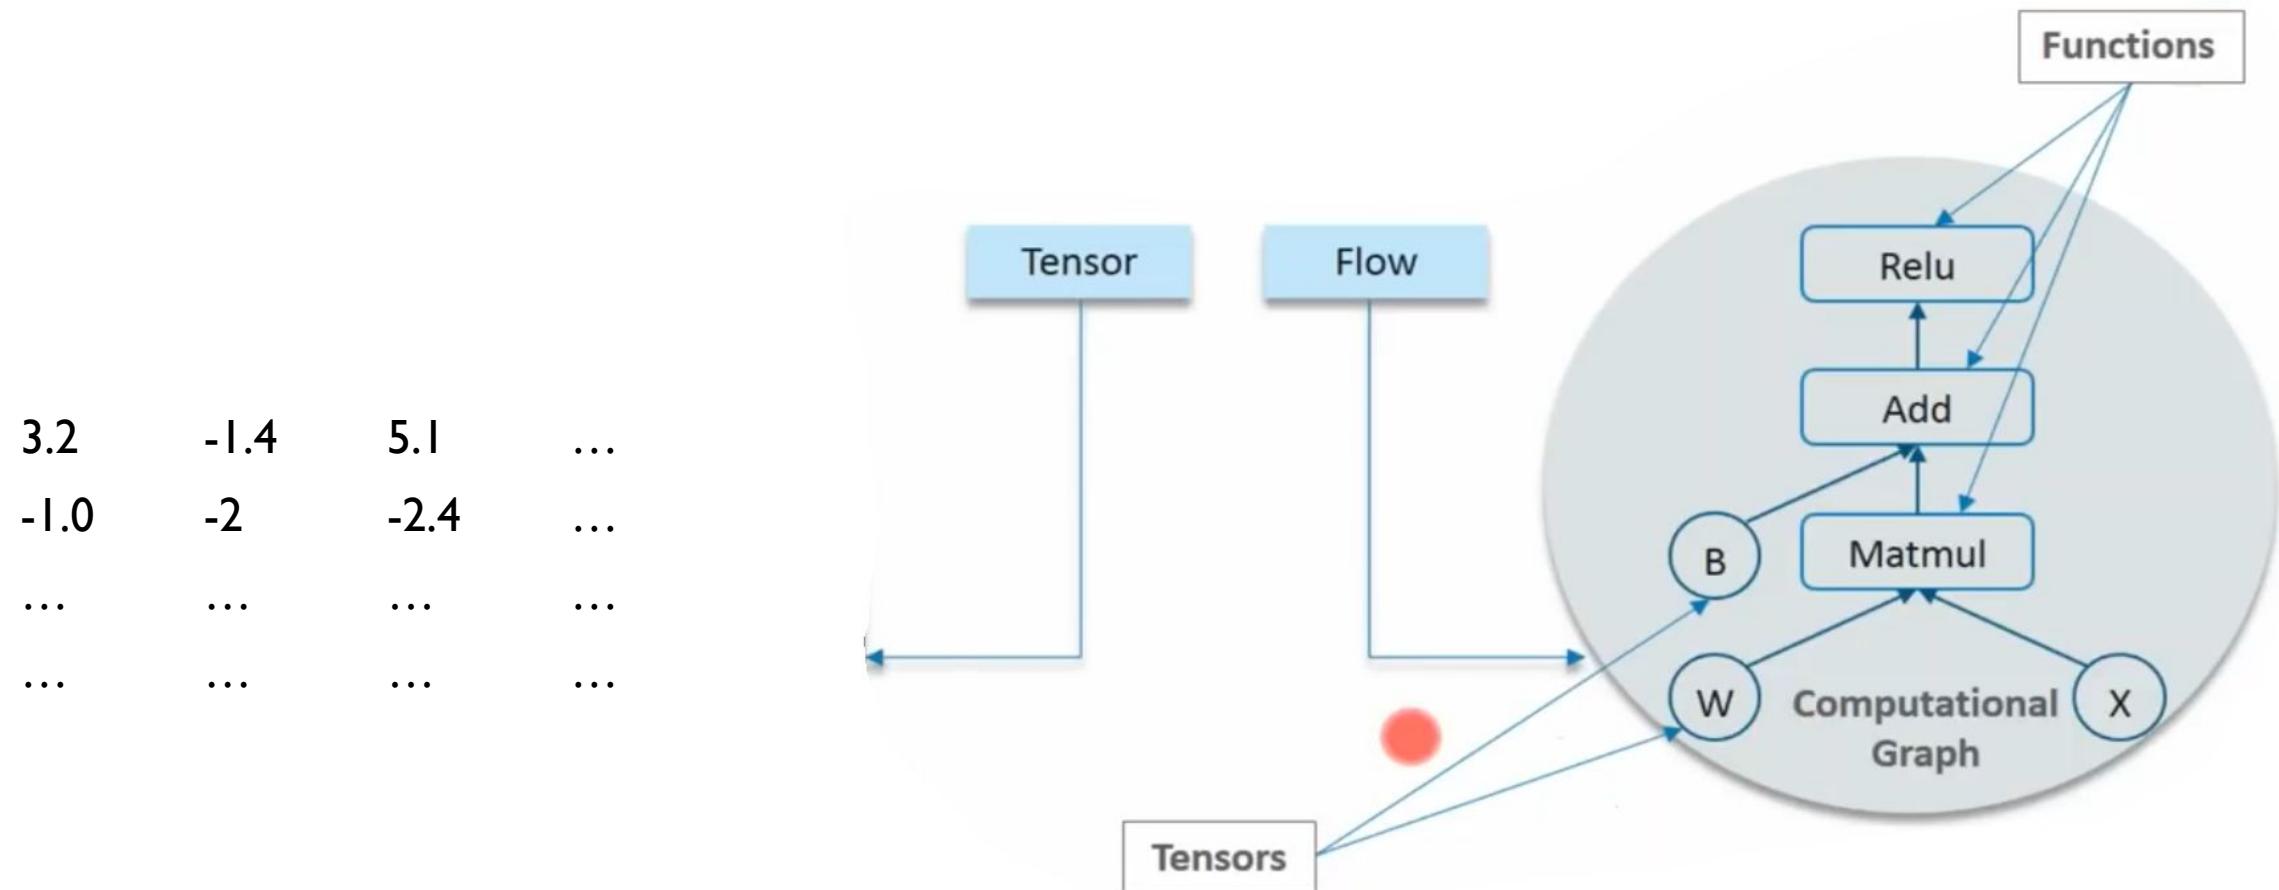
Tensor Data Types

In addition to dimensionality Tensors have different data types as well,you can assign any one of these data types to a tensor.

Data type	Python type	Description
DT_FLOAT	<code>tf.float32</code>	32 bits floating point.
DT_DOUBLE	<code>tf.float64</code>	64 bits floating point.
DT_INT8	<code>tf.int8</code>	8 bits signed integer.
DT_INT16	<code>tf.int16</code>	16 bits signed integer.
DT_INT32	<code>tf.int32</code>	32 bits signed integer.
DT_INT64	<code>tf.int64</code>	64 bits signed integer.
DT_UINT8	<code>tf.uint8</code>	8 bits unsigned integer.
DT_STRING	<code>tf.string</code>	Variable length byte arrays. Each element of a tensor is a byte array.
DT_BOOL	<code>tf.bool</code>	Boolean.

What is Tensor Flow?

- ❑ Tensor Flow is a python library used to implement deep networks.
- ❑ In Tensor Flow, computation is approached as a dataflow graph..



Tensor Flow Code-Basics

Tensor Flow core programs consists of two discrete sections:

Building a
computational graph

Running a
computational graph

A computational graph is a
series of Tensor Flow
operations arranged into a
graph of nodes

Tensor Flow Building And Running A Graph

Building a computational graph

```
Import tensorflow as tf
```

```
node = tf.constant(3.0,tf.float32)  
Node = tf.constant(4.0)  
print(node1,node2)
```

Constant nodes

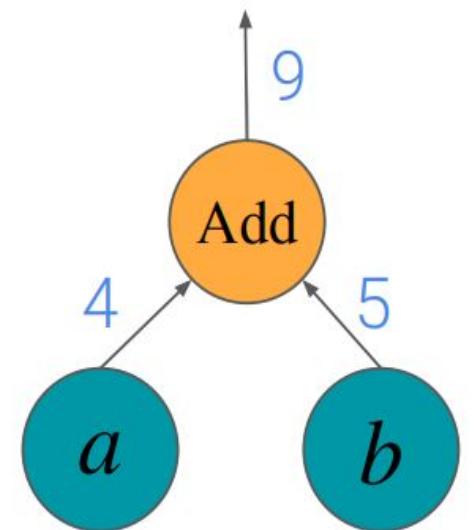
Running a computational graph

```
Sess=tf.Session()  
Print(sess.run([node1,node2]))
```

To actually evaluate the nodes, we must run the computational graph within a **session**. As the session encapsulates the control and state of the tensorflow runtime.

Basic Code Structure

- View functions as computational graphs
- **First** build a computational graph, and Nodes are operators (ops), variables, and constants
- Edges are tensors
 - 0-d is a scalar
 - 1-d is a vector
 - 2-d is a matrix ◦ Etc.
- **TensorFlow = Tensor + Flow = Data + Flow**
- **Second**, use a session to execute operations in the graph



Basic Code Structure - Graphs

- **Constants** are fixed value tensors - not trainable
- **Variables** are tensors initialized in a session – trainable
- **Placeholders** are tensors of values that are unknown during the graph construction, but passed as input during a session
- **Ops** are functions on tensors

Constants

Constant

Placeholders

Variable

One type of node is a constant. It takes no inputs and it outputs a value it stores internally

Import tensorflow as tf

```
node1=tf.constant(3.0,tf.float32)  
node=2.tfconstant(4.0)  
print(node1,node2)
```

Constant nodes

Placeholder

A graph can be parameterized to accept external input, known as placeholders. A placeholder is a promise to provide a value later

Constant

Placeholders

Variable

Constant nodes

Import tensorflow as tf

```
a=tf.placeholder(tf.float32)  
b=tf.placeholder(tf.float32)  
Adder_node = a + b  
Sess=tf.Session()
```

```
Print(sess.run(adder_node,{a:[1,3],b:[2,4]}))
```

Feed values to the
placeholders a and b

Variable

Constant

Placeholders

Variable

To initialize all the variables in a TensorFlow program, you must explicitly call a special operation

Variable
s

To make the model trainable, we need to be able to modify the graph to get new outputs with same input. Variables allow us to add trainable parameters to a graph.

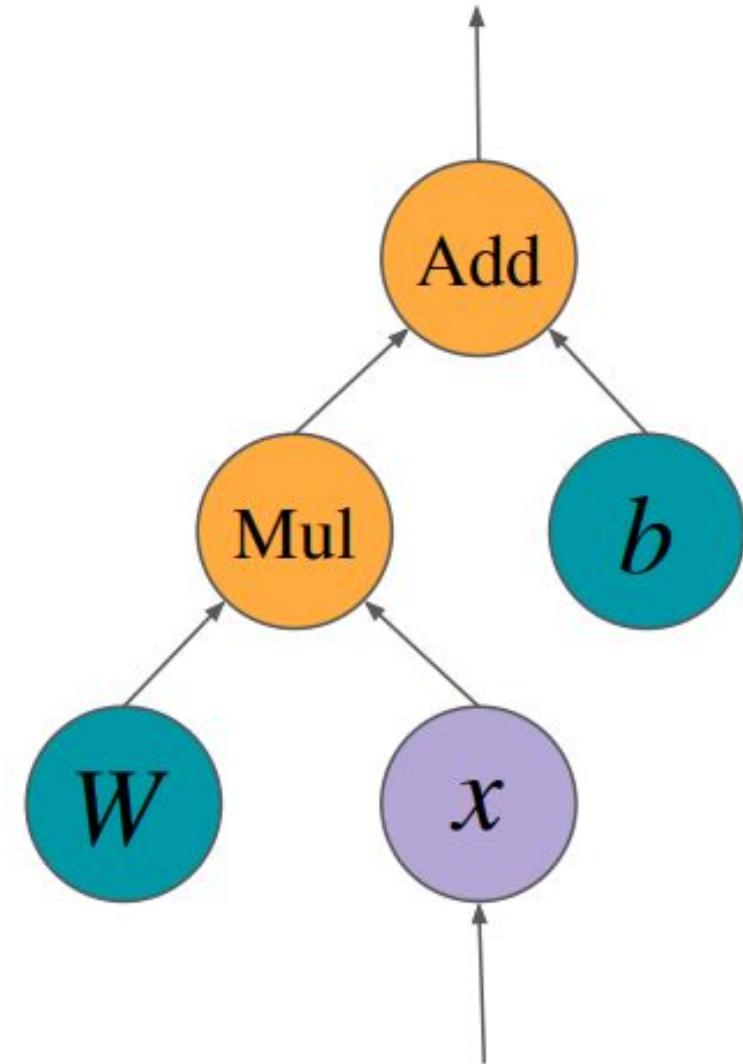
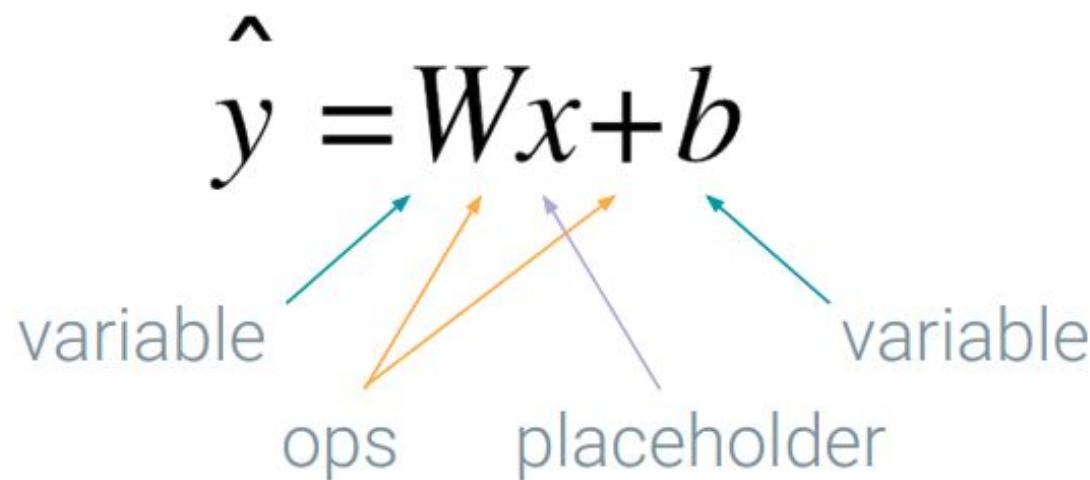
Import tensorflow as tf
`W = tf.Variable([.3], tf.float32)`
`B= tf.Variable([-3], tf.float32)`

`linear_model=W* x + b`
`init = tf.global_variables_initializer()`
`sess = tf.Session()`
`sess.run(init)`
`Print(sess.run(linear_model,{x[1,2,4]}))`

Basic Code Structure - Graphs

$$\hat{y} = Wx + b$$

variable ops placeholder variable



Basic Code Structure - Sessions

- Session is the runtime environment of a graph, where operations are executed, and tensors are evaluated

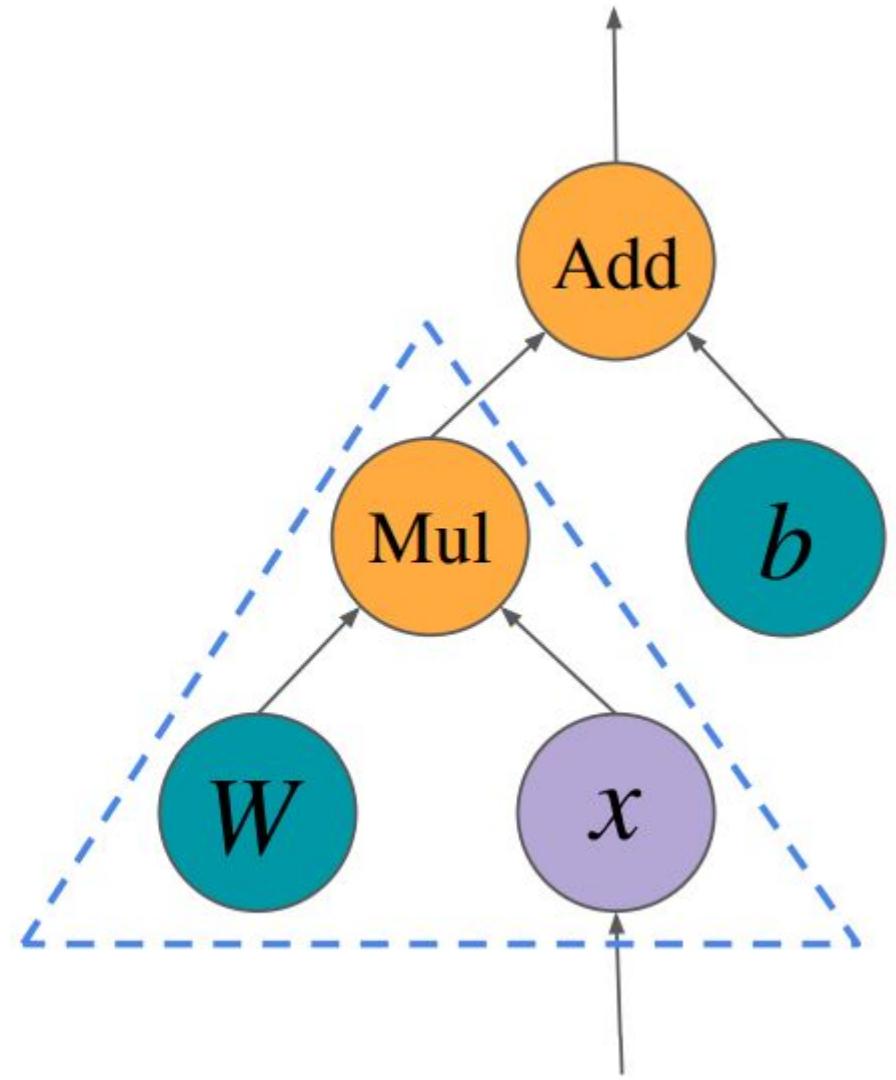
```
>>> import tensorflow as tf  
>>> a = tf.constant(4)  
>>> b = tf.constant(5)  
>>> add_op = tf.add(a, b)  
>>> print(add_op)  
Tensor("Add:0", shape=(), dtype=int32)
```

```
>>> import tensorflow as tf  
>>> a = tf.constant(4)  
>>> b = tf.constant(5)  
>>> add_op = tf.add(a, b)  
>>> with tf.Session() as session:  
...     print(session.run(add_op))  
...  
9
```

- `a.eval()` is equivalent to `session.run(a)`, but in general, “eval” is limited to executions of a single op and ops that returns a value
- Upon op execution, only the subgraph required for calculating its value is evaluated

$$\hat{y} = Wx + b$$

variable
ops placeholder
variable



Abstract Tensor code

The screenshot shows the PyCharm IDE interface with a TensorFlow project open. The code editor displays a Python script named `constant.py` containing the following code:

```
import tensorflow as tf

node1 = tf.constant(5.0, tf.float32)
node2 = tf.constant(4, tf.int16)
node3 = tf.constant(6.0)
node4 = tf.constant(2)

print(node1)
print(node2)
print(node3)
print(node4)
```

The line `node2 = tf.constant(4, tf.int16)` is highlighted with a yellow background. The run output window at the bottom shows the results of the script execution:

```
"C:\Users\Parth Goel\PycharmProjects\TensorFlow\venv\Scripts\python.exe" "C:/Users/Parth Goel/PycharmProj...
Tensor("Const:0", shape=(), dtype=float32)
Tensor("Const_1:0", shape=(), dtype=int16)
Tensor("Const_2:0", shape=(), dtype=float32)
Tensor("Const_3:0", shape=(), dtype=int32)
```

The status bar at the bottom indicates the script finished with exit code 0.

Method 1 - To Execute The Graph

The screenshot shows the PyCharm IDE interface with a TensorFlow project open. The code in `constant.py` defines four constant nodes and prints their values and types.

```
import tensorflow as tf

node1 = tf.constant(5.0, tf.float32)
node2 = tf.constant(4, tf.int16)
node3 = tf.constant(6.0)
node4 = tf.constant(2)

mySess = tf.Session()

print("Node: 1 - ", mySess.run(node1), " Node: 2 - ", mySess.run(node2))
print("Node: 3 - ", mySess.run(node3), " Node: 4 - ", mySess.run(node4))

print(node1, node2)
```

The output window shows the results of the session runs:

```
"C:\Users\Parth Goel\PycharmProjects\TensorFlow\venv\Scripts\python.exe" "C:/Users/Parth Goel/PycharmProj...
Node: 1 - 5.0      Node: 2 - 4
Node: 3 - 6.0      Node: 4 - 2
Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=int16)
Tensor("Const_2:0", shape=(), dtype=float32) Tensor("Const_3:0", shape=(), dtype=int32)
```

Method 2 - To Execute The Graph

The screenshot shows the PyCharm IDE interface with the following details:

- Title Bar:** TensorFlow [C:\Users\Parth Goel\PycharmProjects\TensorFlow] - ...\\Code\\constant.py [TensorFlow] - PyCharm
- Menu Bar:** File Edit View Navigate Code Refactor Run Tools VCS Window Help
- Toolbars:** standard toolbar icons
- Project Explorer:** TensorFlow project structure with Code, constant.py, venv library root, External Libraries, and Scratches and Consoles.
- Code Editor:** constant.py file containing Python code to create four TensorFlow constants and print their values. The code uses tf.Session() to run operations.

```
import tensorflow as tf
node1 = tf.constant(5.0, tf.float32)
node2 = tf.constant(4, tf.int16)
node3 = tf.constant(6.0)
node4 = tf.constant(2)

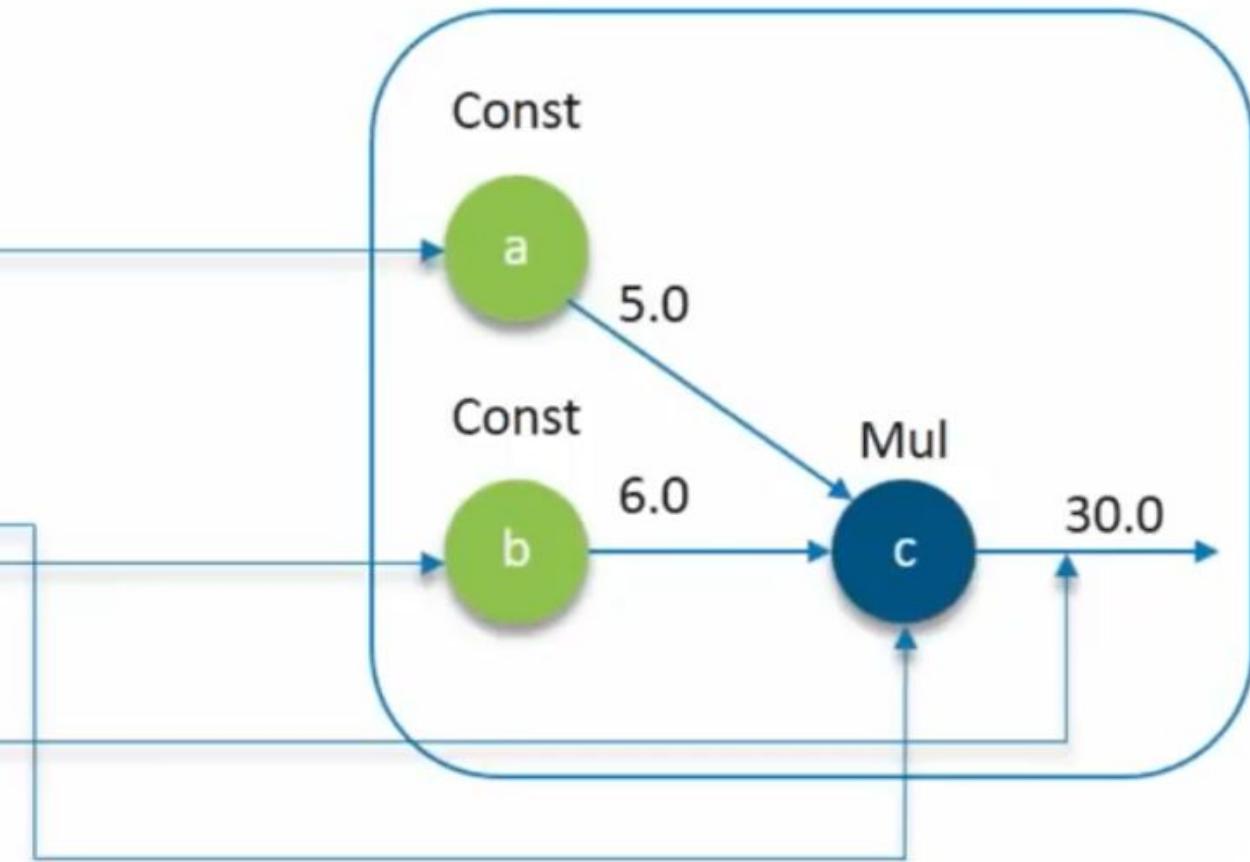
with tf.Session() as mySess:
    print("Node: 1 - ",mySess.run(node1), " Node: 2 - ",mySess.run(node2))
    print("Node: 3 - ",mySess.run(node3), " Node: 4 - ",mySess.run(node4))

    print(node1,node2)
    print(node3,node4)
```

- Run Tab:** Shows the command being run: "C:\\Users\\Parth Goel\\PycharmProjects\\TensorFlow\\venv\\Scripts\\python.exe" "C:/Users/Parth Goel/PycharmProj...".
- Output Tab:** Displays the execution results:
 - Printed values: Node: 1 - 5.0 Node: 2 - 4
Node: 3 - 6.0 Node: 4 - 2
 - Tensor objects:
 - Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=int16)
 - Tensor("Const_2:0", shape=(), dtype=float32) Tensor("Const_3:0", shape=(), dtype=int32)

Tensorflow Example

```
import tensorflow as tf  
  
# Build a graph  
a = tf.constant(5.0)  
b = tf.constant(6.0)  
  
c = a * b  
  
# Launch the graph in a session  
sess = tf.Session()  
  
# Evaluate the tensor 'C'  
print(sess.run(c))
```



Arithmetic Operations Example

The screenshot shows the PyCharm IDE interface with the following details:

- Title Bar:** TensorFlow [C:\Users\Parth Goel\PycharmProjects\TensorFlow] - ...\\Code\\arith_op.py [TensorFlow] - PyCharm
- Menu Bar:** File Edit View Navigate Code Refactor Run Tools VCS Window Help
- Toolbars:** TensorFlow, Code, arith_op.py
- Project Explorer:** TensorFlow project structure with files arith_op.py and constant.py.
- Code Editor:** Content of arith_op.py:

```
import tensorflow as tf

number1 = tf.constant(6)
number2 = tf.constant(5)

add = number1 + number2
sub = number1 - number2
mul = number1 * number2
div = number1 / number2

with tf.Session() as mySess:
    print("ADD - ",mySess.run(add), " SUB - ",mySess.run(sub))
    print("MUL - ",mySess.run(mul), " DIV - ",mySess.run(div))
```
- Run Tab:** Run configuration for arith_op.
- Output Tab:** Shows the execution results:

```
Process finished with exit code 0
```

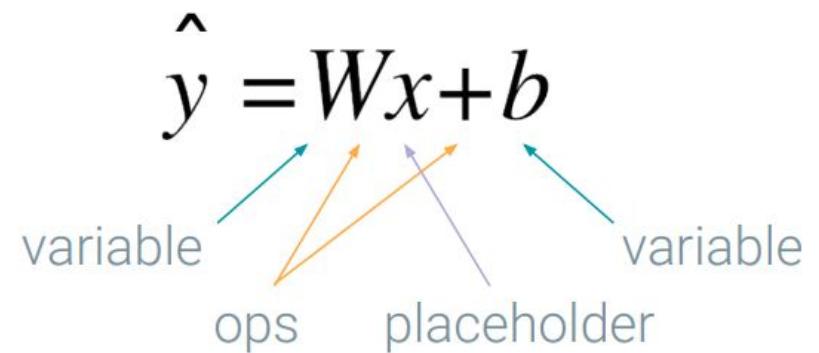
ADD - 11	SUB - 1
MUL - 30	DIV - 1.2
- Bottom Status Bar:** 7:24 CRLF UTF-8, 9:31 PM

Simple linear Model

```
import tensorflow as tf
W= tf.Variable([0.5], tf.float32)
b = tf.Variable([-0.5], tf.float32)
X= tf.placeholder(tf.float32)

linear_model = W * X + b
init= tf.global_variables_initializer()

sess= tf.Session()
sess.run(init)
print(sess.run(linear_model, {X:[1,2,3,4]}))
```

$$\hat{y} = Wx + b$$


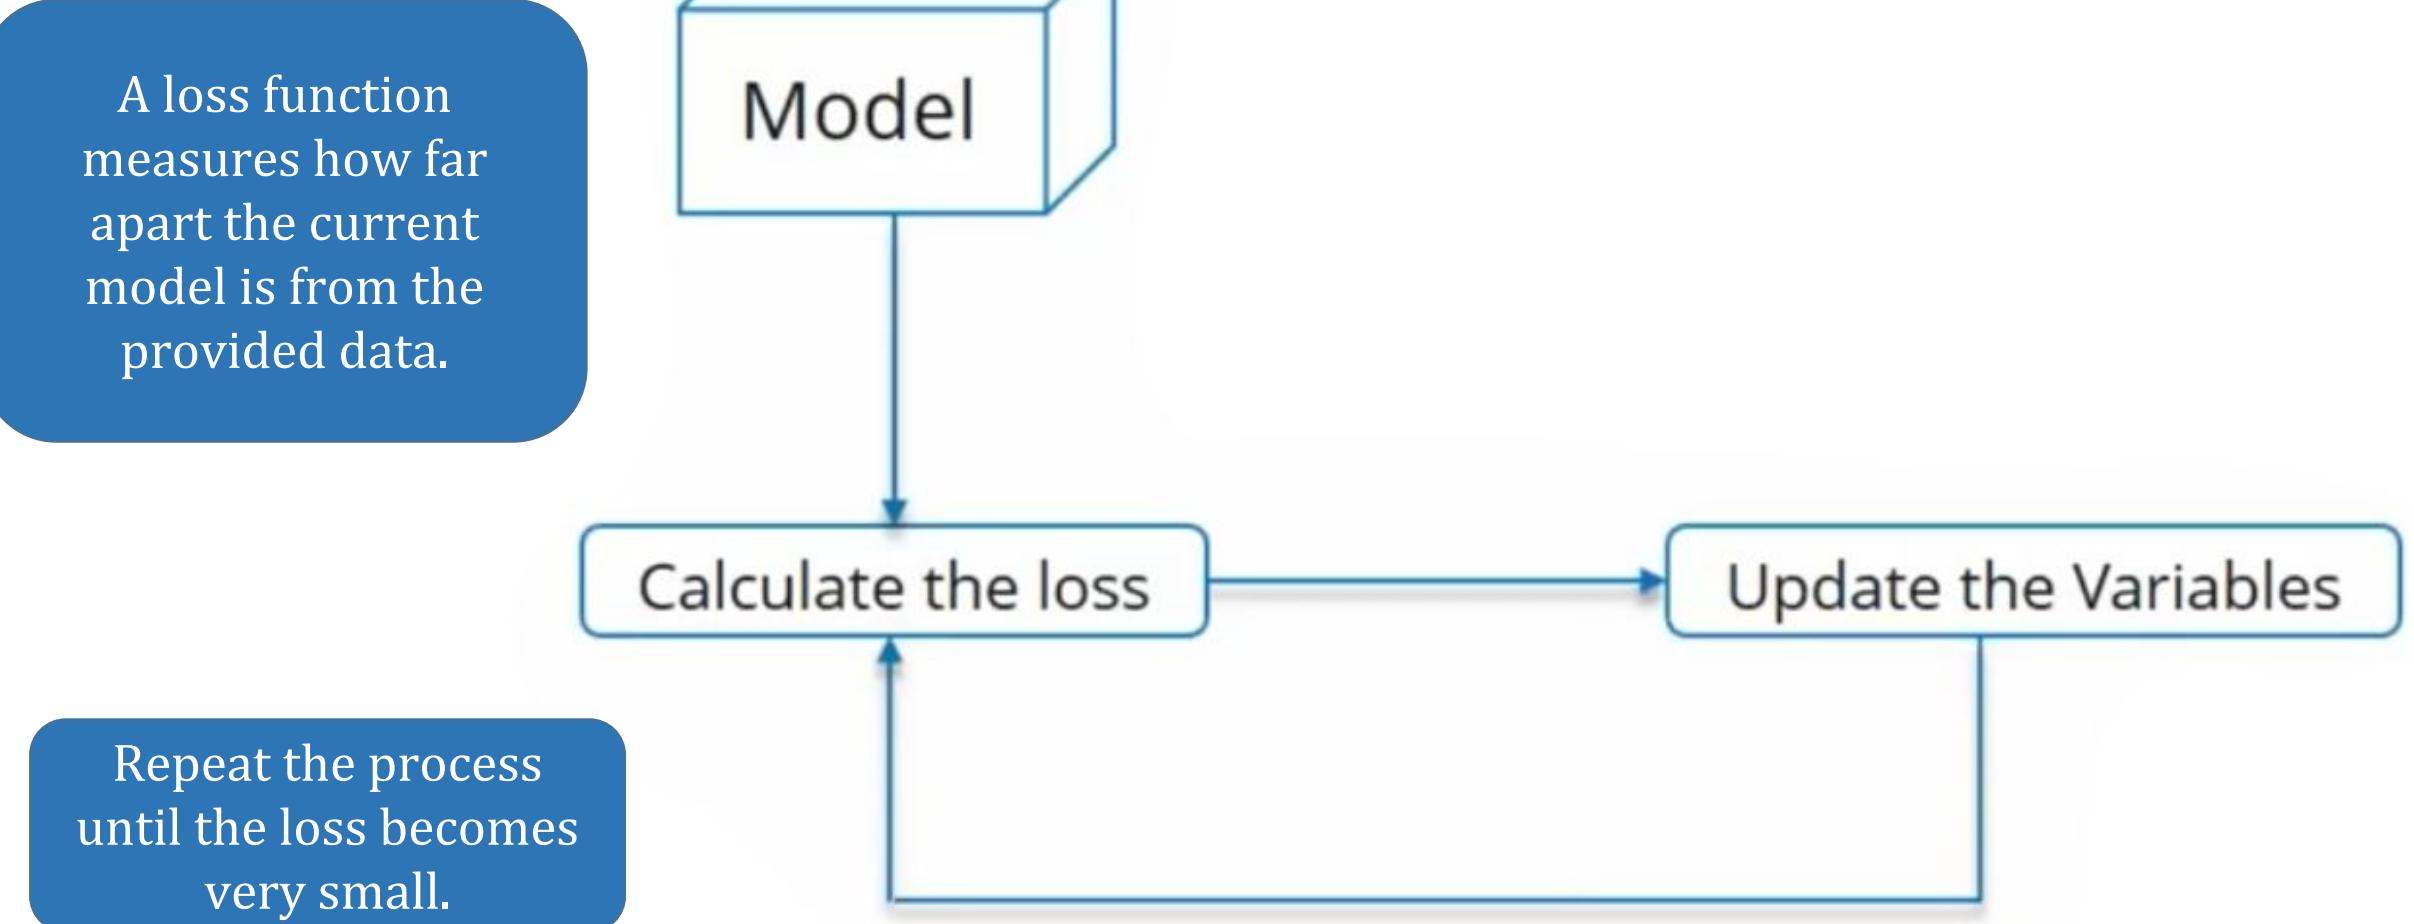
variable
ops placeholder
variable

```
1 import tensorflow as tf
2 W= tf.Variable([0.5], tf.float32)
3 b = tf.Variable([-0.5], tf.float32)
4 X= tf.placeholder(tf.float32)
5
6 linear_model = W * X + b
7
8 init= tf.global_variables_initializer()
9
10 sess= tf.Session()
11 sess.run(init)
12 print(sess.run(linear_model, {X:[1,2,3,4]}))
```

```
"C:\\Users\\Parth Goel\\PycharmProjects\\TensorFlow\\venv\\Scripts\\python.exe" "C:\\Users\\Parth Goel\\PycharmProj...  
[0. 0.5 1. 1.5]
```

Process finished with exit code 0

How To Increase The efficiency Of The Model?



Calculating The Loss

In order to understand how good the model is, we should know the loss/error

```
y=tf.placeholder(tf.float32)  
  
Sqrard_deltas =  
tf.square(linear_model -y)  
loss=tf.reduce_sum(squared_deltas)  
print(sess.run(loss,{x:[1,2,3,4]  
,y:[0,-1,-2,-3]}))
```

To evaluate the model on training data, we need a y i.e a placeholder to provide the desired values, and we need to write a loss function

We'll use a standard loss model for linear regression. $(\text{linear_model}-y)$ creates a vector where each element is the corresponding example's error delta.

`tf.square` is used to square that error.
`tf.reduce_sum` is used to sum all the squared error.

TensorFlow [C:\Users\Parth Goel\PycharmProjects\TensorFlow] - ...\\Code\\LM.py [TensorFlow] - PyCharm

File Edit View Navigate Code Refactor Run Tools VCS Window Help

TensorFlow > Code > LM.py

Project constant.py arith_op.py LM.py

1: Project TensorFlow C:\Users\Parth Goel\PycharmProjects\TensorFlow\Code
Code
constant.py
arith_op.py
LMI.py
LM.py
venv library root
External Libraries
Scratches and Consoles

2: Structure

3: Favorites

```
1 import tensorflow as tf
2 W = tf.Variable([0.5], tf.float32)
3 b = tf.Variable([-0.5], tf.float32)
4 X = tf.placeholder(tf.float32)
5
6 linear_model = W * X + b
7
8 y = tf.placeholder(tf.float32)
9
10 squared_deltas = tf.square(linear_model - y)
11 loss = tf.reduce_sum(squared_deltas)
12
13 init = tf.global_variables_initializer()
14
15 sess = tf.Session()
16 sess.run(init)
17 print(sess.run(loss, {X: [1, 2, 3, 4], y: [0, -1, -2, -3]}))
```

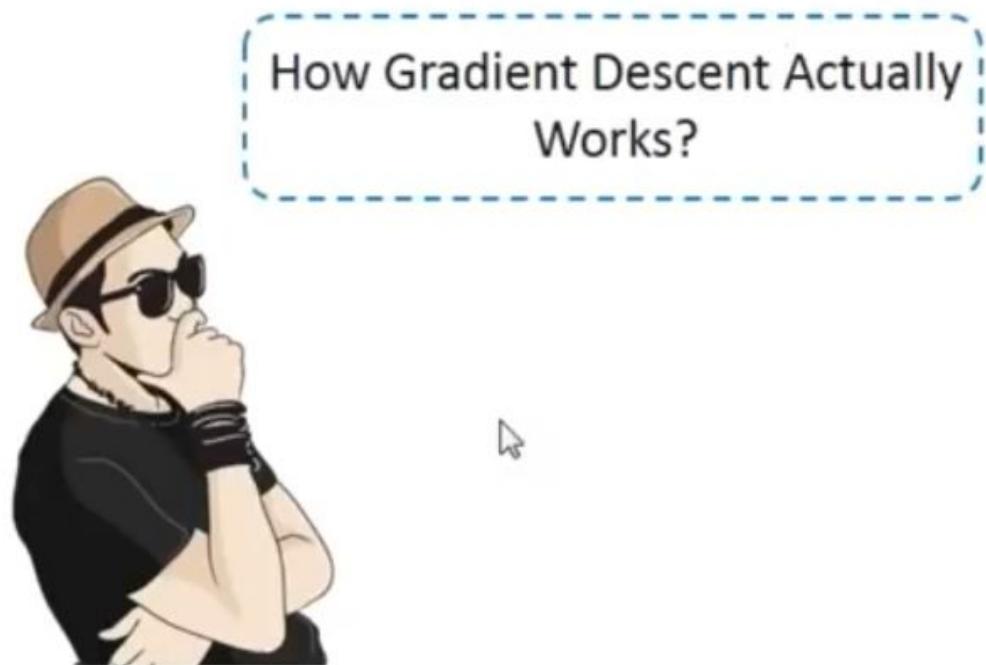
Run: LM

C:\Users\Parth Goel\PycharmProjects\TensorFlow\venv\Scripts\python.exe "C:/Users/Parth Goel/PycharmProj...
31.5

4: Run 6: TODO Terminal Python Console Event Log

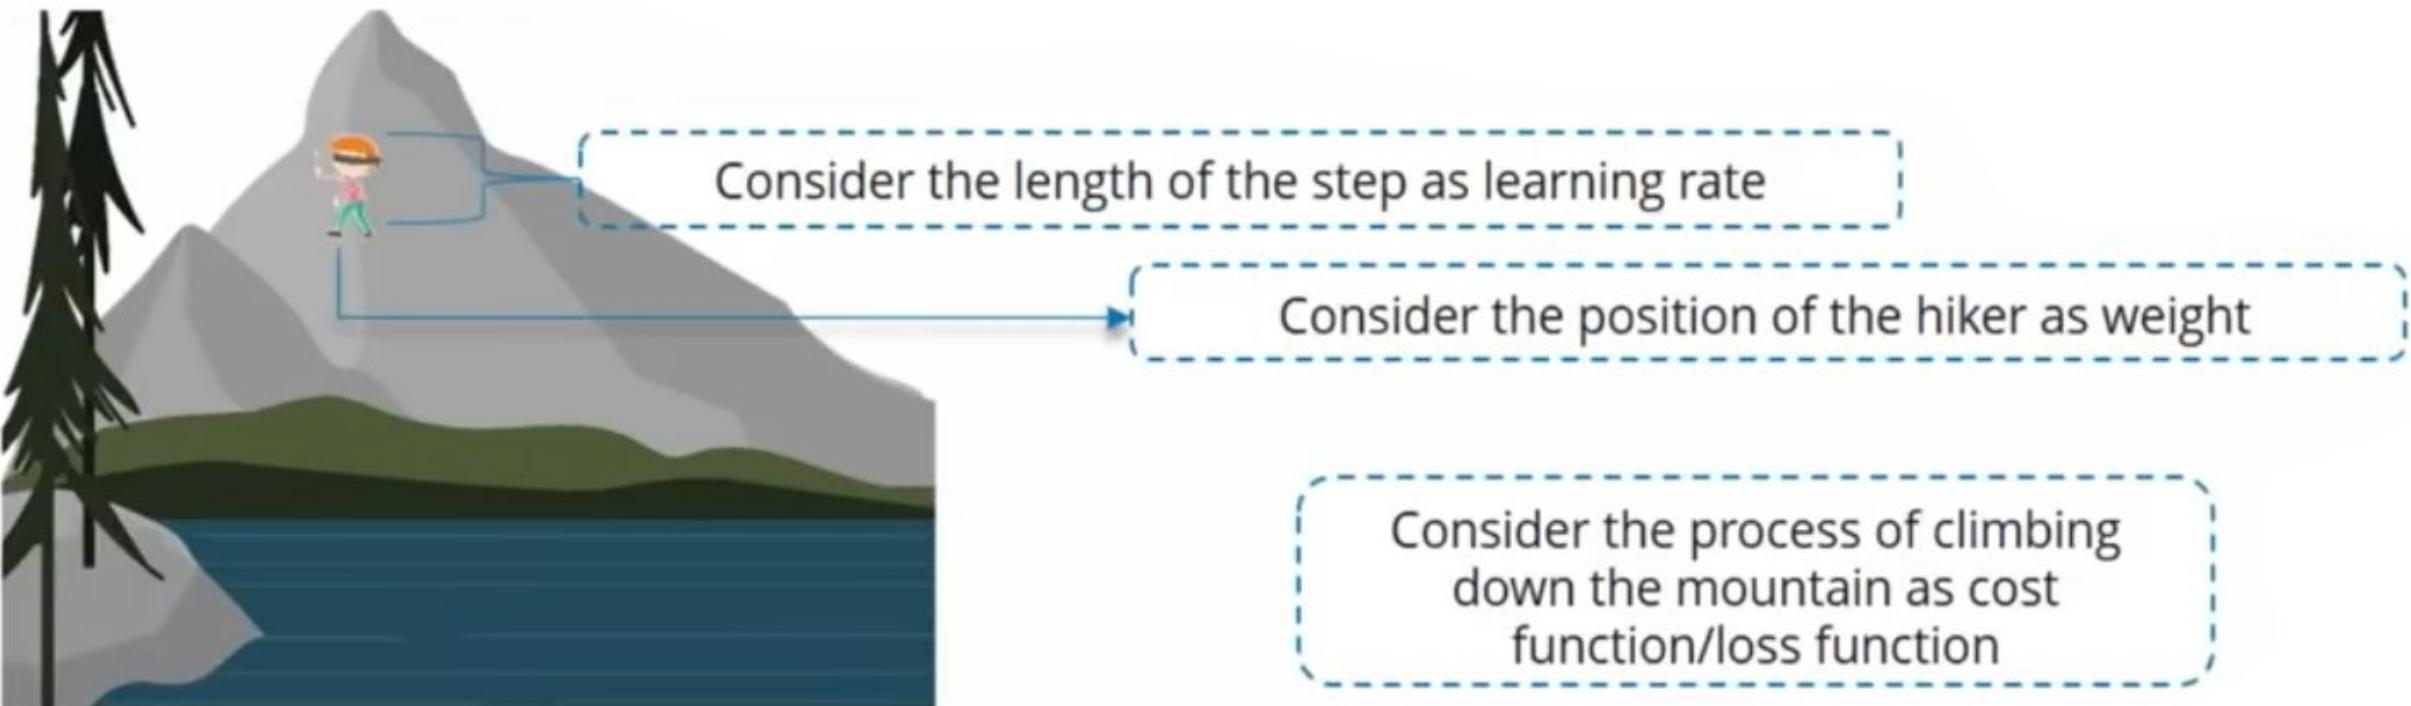
Reducing The Loss

Optimizer modifies each variable according to the magnitude of the derivative of a loss with respect to that variable. Here we will use Gradient descent optimizer.



Reducing The Loss

- Suppose you are at the top of a mountain and you have to reach a lake which is at the lowest point of the mountain(aka valley).
- A twist is that you are blindfolded and you have zero visibility to see where you are headed. So, what approach will take to reach the lake?



Gradient Descent

The weights are updated incrementally after each epoch. The cost function $j(\cdot)$, the sum of squared errors(SSE), can be written as:

$$J(\mathbf{w}) = \frac{1}{2} \sum_i (\text{target}^{(i)} - \text{output}^{(i)})^2$$

The magnitude and direction of the weight update is computed by taking a step in the opposite direction of the cost gradient.

$$\Delta w_j = -\eta \frac{\partial J}{\partial w_j},$$

The weights are then updated after each epoch via the following update rule:

$$\mathbf{w} := \mathbf{w} + \Delta \mathbf{w},$$

Here, $\Delta \mathbf{w}$ is a vector that contains the weight updates of each weight coefficient w , which are computed as follows:

$$\Delta w_j = -\eta \frac{\partial J}{\partial w_j} = -\eta \sum_i (\text{target}^{(i)} - \text{output}^{(i)})(-x_j^{(i)}) = \eta \sum_i (\text{target}^{(i)} - \text{output}^{(i)})x_j^{(i)}.$$

Reducing The Loss

Suppose, we want to find the best parameters (W) for our learning algorithm. We can apply the same analogy and find the best possible values for that parameter, Consider the example below:

```
Optimizer = tf.train.GradientDescentOptimizer(0.01)
train = optimizer.minimize(loss)

sess.run(init)
for i in range(1000):
    sess.run(train, {x:[1,2,3,4],y:[0,-1,-2,-3] })
print(sess.run([W,b]))
```

```
1 import tensorflow as tf
2 W = tf.Variable([0.5], tf.float32)
3 b = tf.Variable([-0.5], tf.float32)
4 X = tf.placeholder(tf.float32)
5
6 linear_model = W * X + b
7 y = tf.placeholder(tf.float32)
8
9 squared_deltas = tf.square(linear_model - y) #Loss
10 loss = tf.reduce_sum(squared_deltas)
11
12 optimizer = tf.train.GradientDescentOptimizer(0.01) #Train
13 train = optimizer.minimize(loss)
14
15 init = tf.global_variables_initializer()
16 sess = tf.Session()
17 sess.run(init)
18
19 for i in range(1000):
20     sess.run(train, {X:[1,2,3,4], y:[0,-1,-2,-3] })
21 print(sess.run([W,b]))
```

```
"C:\Users\Parth Goel\PycharmProjects\TensorFlow\venv\Scripts\python.exe" "C:/Users/Parth Goel/PycharmProjects/TensorFlow/Code/
```

Graph Visualization

- For visualizing Tensorflow graphs, we use Tensor Board.
- The first argument when creating the file Writer is an output **directory name**, which will be created if it doesn't exist.

```
File_writer = tf.summary.FileWriter("log_simple_graph", mySess.graph)
```

truediv



```
tensoboard -logdir = "path_to_the_graph"
```

TensorBoard runs as a local web app, on port 6006.
(this is default port.)

How To Run Tensorboard

- In PyCharm
 - Go to VIEW -> Tool Windows -> Terminal
 - >>>python venv\Lib\site-packages\tensorboard\main.py --logdir=<dir.name>
- (Make sure that python3.5 path is set in environment variable.)

TensorBoard

GRAPHS

INACTIVE



Search nodes. Regexes sup...



Fit to screen



Download PNG

Run

(1)

Session runs (0)

Upload

Choose File

 Trace inputs

▼ Close legend.

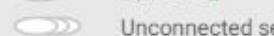
Graph (* = expandable)



Namespace*?



OpNode?



Unconnected series*?



Connected series*?



Constant?



Summary?



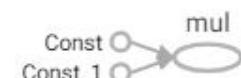
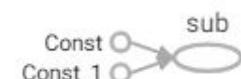
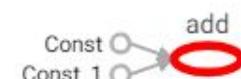
Dataflow edge?



Control dependency edge?



Reference edge?



add

Operation: Add

Attributes (1)

T {"type":"DT_INT32"}

Inputs (2)

- Const
- Const_1

Outputs (0)

Add to main graph



Main Graph

Auxiliary Nodes

Search nodes. Regexes sup...

Fit to screen

Download PNG

Run
(1)Session
runs (0)Upload Trace inputs

Close legend.

Graph (* = expandable)

Namespace?

OpNode?

Unconnected series?

Connected series?

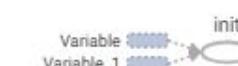
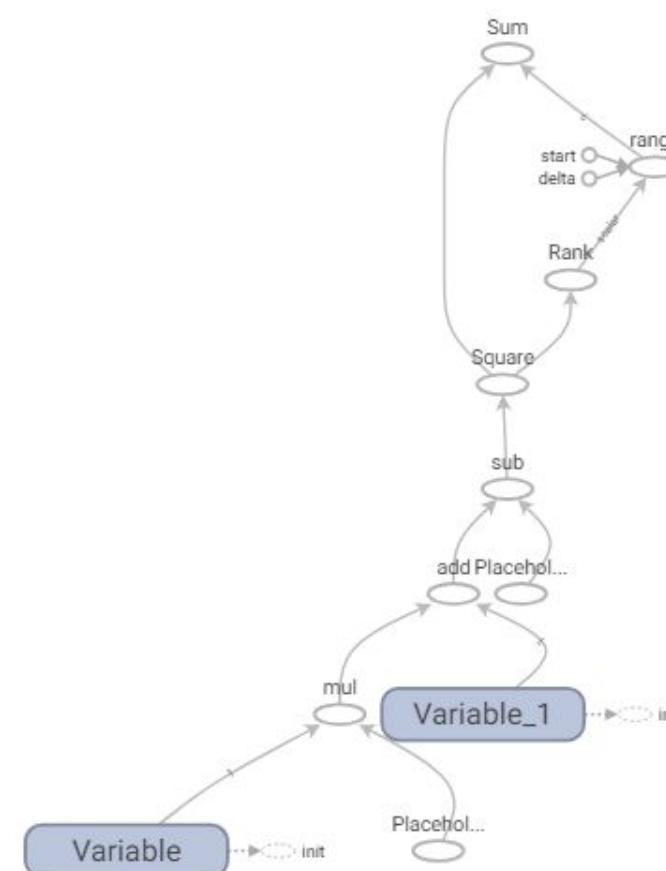
Constant?

Summary?

Dataflow edge?

Control dependency edge?

Reference edge?



```

import tensorflow as tf
W= tf.Variable([0.5], tf.float32)
b = tf.Variable([-0.5], tf.float32)
X= tf.placeholder(tf.float32)
  
```

```

linear_model = W * X + b
init= tf.global_variables_initializer()
  
```

```

sess= tf.Session()
sess.run(init)
File_Writer =
tf.summary.FileWriter('graph_LM',sess.graph)
  
```

```

print(sess.run(linear_model,{X:[1,2,3,4]}))
  
```

TensorFlow Cheat Sheet

About

TensorFlow

TensorFlow™ is an open source software library for numerical computation using data flow graphs. TensorFlow was originally developed for the purposes of conducting machine learning and deep neural networks research, but the system is general enough to be applicable in a wide variety of other domains as well.

Skflow

Scikit Flow provides a set of high level model classes that you can use to easily integrate with your existing Scikit-learn pipeline code. Scikit Flow is a simplified interface for TensorFlow, to get people started on predictive analytics and data mining. Scikit Flow has been merged into TensorFlow since version 0.8 and now called TensorFlow Learn.

Keras

Keras is a minimalist, highly modular neural networks library, written in Python and capable of running on top of either TensorFlow or Theano

Installation

How to install new package in Python:

```
pip install <package-name>
```

Example: `pip install requests`

How to install tensorflow?

```
device = cpu/gpu
```

```
python_version = cp27/cp34
```

```
sudo pip install
```

```
https://storage.googleapis.com/tensorflow/linux/$device/tensorflow-0.8.0-$python_version-none-linux_x86_64.whl
```

How to install Skflow

```
pip install sklearn
```

How to install Keras

```
pip install keras
```

update `~/.keras/keras.json` - replace "theano" by "tensorflow"

Helpers

Python helper

Important functions

`type(object)`

Get object type

`help(object)`

Get help for object (list of available methods, attributes, signatures and so on)

`dir(object)`

Get list of object attributes (fields, functions)

`str(object)`

Transform an object to string

`object?`

Shows documentations about the object

`globals()`

Return the dictionary containing the current scope's global variables.

`locals()`

Update and return a dictionary containing the current scope's local variables.

`id(object)`

Return the identity of an object. This is guaranteed to be unique among simultaneously existing objects.

`import __builtin__ dir(__builtin__)`

Other built-in functions

TensorFlow

Main classes

`tf.Graph()`

`tf.Operation()`

`tf.Tensor()`

`tf.Session()`

Some useful functions

`tf.get_default_session()`

`tf.get_default_graph()`

`tf.reset_default_graph()`

`ops.reset_default_graph()`

`tf.device("/cpu:0")`

`tf.name_scope(value)`

`tf.convert_to_tensor(value)`

TensorFlow Optimizers

`GradientDescentOptimizer`

`AdadeltaOptimizer`

`AdagradOptimizer`

`MomentumOptimizer`

`AdamOptimizer`

`FtrlOptimizer`

`RMSPropOptimizer`

Reduction

reduce_sum
reduce_prod
reduce_min
reduce_max
reduce_mean
reduce_all
reduce_any
accumulate_n

Activation functions

tf.nn?

relu

relu6

elu

softplus

softsign

dropout

bias_add

sigmoid

tanh

sigmoid_cross_entropy_with_logits

softmax

log_softmax

softmax_cross_entropy_with_logits

sparse_softmax_cross_entropy_with_logits

weighted_cross_entropy_with_logits

etc.

Skflow

Main classes

TensorFlowClassifier
TensorFlowRegressor
TensorFlowDNNClassifier
TensorFlowDNNRegressor
TensorFlowLinearClassifier
TensorFlowLinearRegressor
TensorFlowRNNClassifier
TensorFlowRNNRegressor

TensorFlowEstimator

Each classifier and regressor have following fields

n_classes=0 (Regressor), n_classes are expected to be input (Classifiers)

batch_size=32,

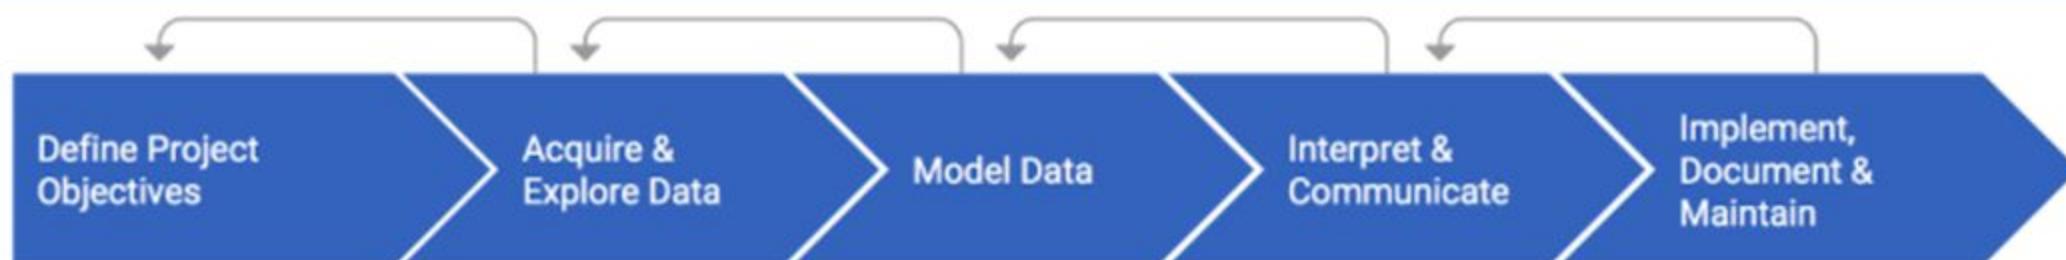
steps=200, // except

TensorFlowRNNClassifier - there is 50

optimizer='Adagrad',

learning_rate=0.1,

The Machine Learning Life Cycle



1. Define Project Objectives

- Specify business problem
- Acquire subject matter expertise
- Define unit of analysis and prediction target
- Prioritize modeling criteria
- Consider risks and success criteria
- Decide whether to continue

2. Acquire & Explore Data

- Find appropriate data
- Merge data into single table
- Conduct exploratory data analysis
- Find and remove any target leakage
- Feature engineering

3. Model Data

- Variable selection
- Build candidate models
- Model validation and selection

4. Interpret & Communicate

- Interpret model
- Communicate model insights

5. Implement, Document & Maintain

- Set up batch or API prediction system
- Document modeling process for reproducibility
- Create model monitoring and maintenance plan

Parameter and Hyperparameter

Parameter - Weight and Bias

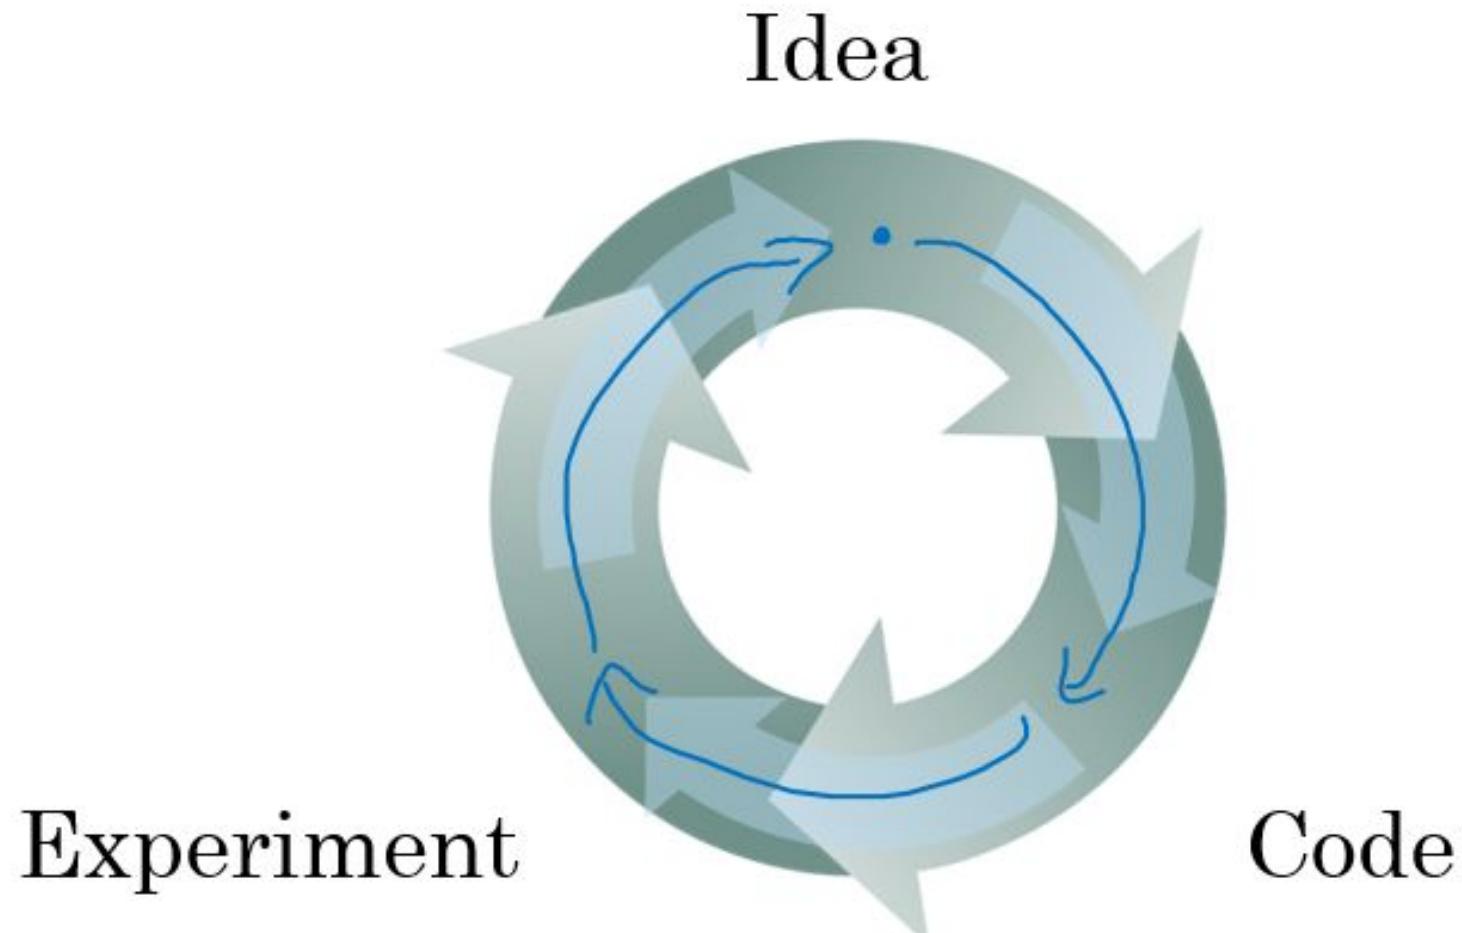
Hyperparameter (parameters that control the algorithm):

- 1) No of hidden layers - Try and error
- 2) No of neurons in each hidden layer - Try and error
- 3) Learning rate. (The most important parameter) – 0.1, 0.01, 0.001, 0.0001, .00001
- 4) Activation functions – Sigmoid, Relu, Leaky Relu, Tanh
- 5) Number of iteration – Try and error
- 6) Epoch – Try and error
- 7) Batch size – 4, 8, 16, 32, 64, 128... (Power of 2)
- 8) Regularization – Dropout, L1, L2
- 9) Regularization Rate and Dropout lambda
- 10) Normalizing inputs – Min-max, mean, Z-Score
- 11) Weights and Bias Initialization – Zero, Random, He
- 12) optimization algorithm – SGD, ADAM, RMSProp
- 13) Learning rate decay
- 14) momentum – 0.9

Iterations, Epoch and Batch Size

- **batch size** = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.
- number of **iterations** = number of passes, each pass using [batch size] number of examples. To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).
- one **epoch** = one forward pass and one backward pass of *all* the training examples.
- **Example:** if you have 5000 training examples, and your batch size is 32, then it will take 157 iterations to complete 1 epoch.
- **$5000/32 = 156.25 = 157$ iterations ($156 \times 32 + 1 \times 8$)**
- **If you set epoch = 20, $20 \times 157 = 3140$ iterations**

Empirical Process



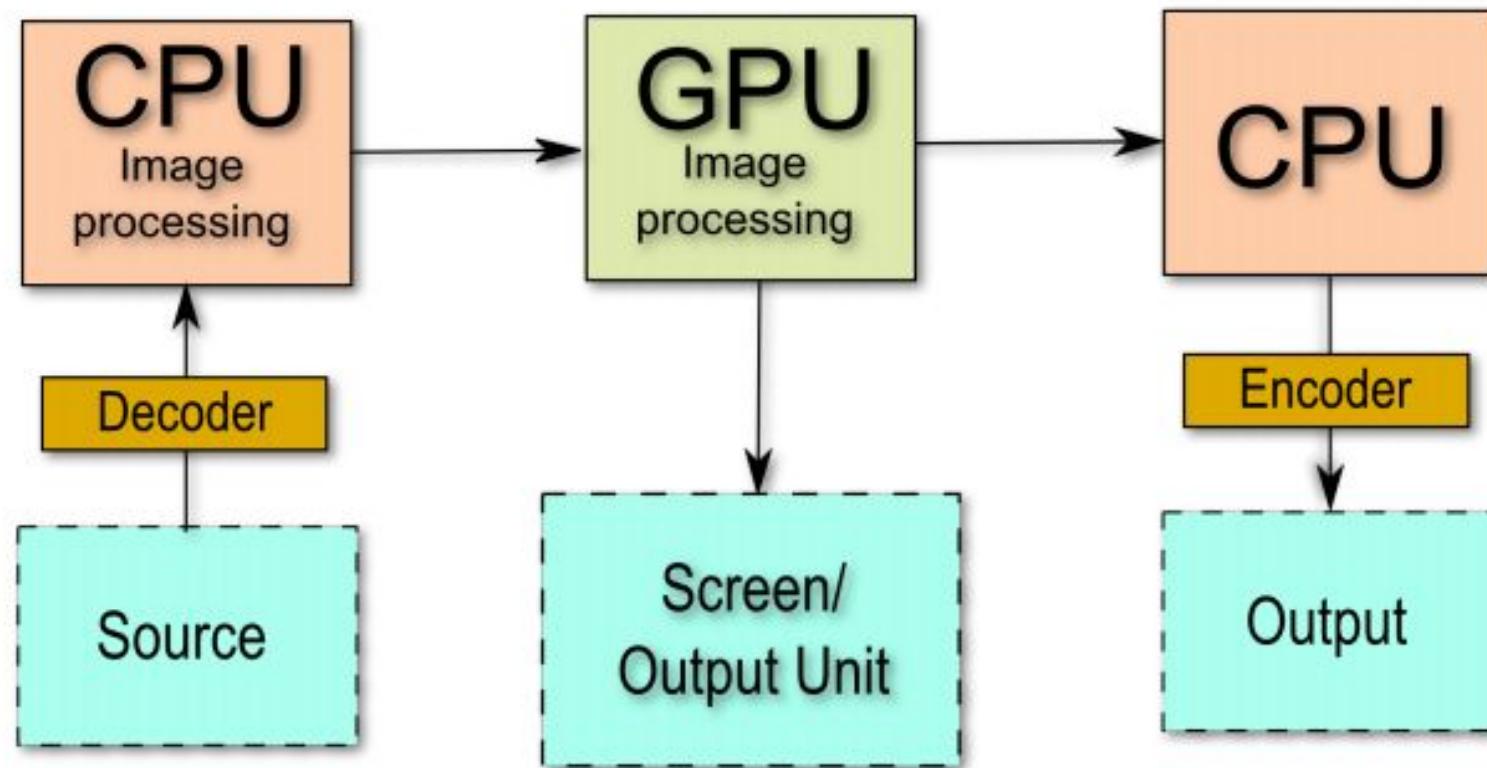
Practical issues with neural networks

- Data augmentation
- Data normalization
- Architecture optimization
- Loss function
- Weight initialization
- Learning rate and its decay schedule
- Overfitting: regularization, early stopping
- Momentum

Research issues in Deep Learning

- Huge amount of data
 - Overfitting in Neural Networks
 - Hyperparameter Optimization
- High-performance Hardware Requirement
- Essentially – It is a Blackbox
- Lack of Flexibility, Scalability and Multitasking
- Generative Models – learn and then generate new images
- Video Processing and Analytics
- Deep Learning in Embedded System
- Generalization of algorithms – (Partial Solution: Transfer Learning)
- CPU to GPU Data Transfer Time

A typical workflow of images/video processing



Deep Learning Library



DEEPMLEARNING4J



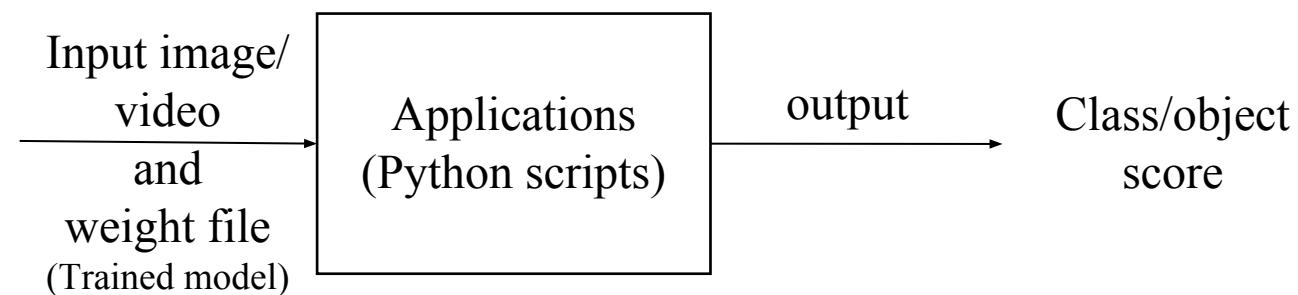
Deep Learning Cloud Services



Comparison of Deep Learning Library

Software	Creator	Software license	Open source	Platform	Written in	Interface	CUDA support
TensorFlow	Google Brainteam	Apache 2.0	Yes	Linux, Mac OS, Windows	C++,Python	Python, (C/C++public API only for executing graphs)	Yes
Keras	François Chollet	MIT license	Yes	Linux, Mac OS, Windows	Python	Python	Yes
Deeplearning4j	Adam Gibson	Apache 2.0	Yes	Linux, Mac OS, Windows, Android (Cross-platform)	C, C++	Java, Scala	Yes
Caffe	Yangqing Jia	BSD License	Yes	Linux, Mac OS, Windows	C++	Python, Matlab	Yes
Theano	Université de Montréal	BSD License	Yes	Cross-platform	Python	Python	Yes
Torch	Ronan Collobert,Kavukcuoglu, Clement	BSD License	Yes	Linux, Mac OS, Windows, Android, iOS	C, Lua	Lua, C, utility library for C++/OpenCL	Yes
Microsoft Cognitive Toolkit - CNTK	Microsoft Research	MIT license	Yes	Windows, Linux	C++	Python, C++, Command line, BrainScript(.NET on roadmap)	Yes

Schematic Diagram



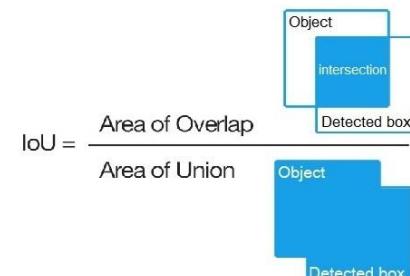
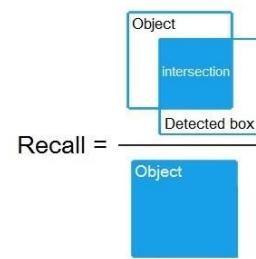
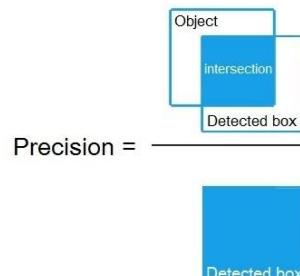
Evaluation Parameters

		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

- Accuracy – Confusion matrix
- Precision
- Recall
- F1 Score (try to maximize - to create a balanced classification model)
- mAP – IOU (For Object Detection)
- Loss/Cost – Minimize the difference between actual and predicted
- Entropy

$$\text{recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

$$\text{precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$



$$F_1 = 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$



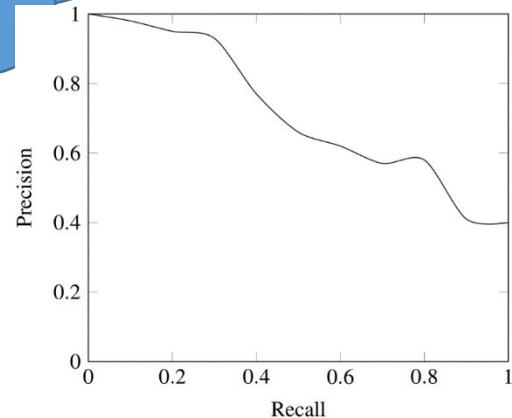
Dear, Do you remember all wedding anniversary surprises from me?

They are celebrating 6th wedding anniversary.

If you can recall all **5** gifts correctly, then, your recall ratio is **1.0 (100%)**. If you can recall 4 events correctly, your recall ratio is **0.8 (80%)**.

After some days, again she asked same question.....

Most
Dangerous
Question!!!



For example, you answers 15 times, 10 times are correct and 5 times are wrong. This means you can recall all events but it's not so precise.

That means: (15 answers: 10 correct answers, 5 wrong answers), you get **100% recall** but your precision is only **66.67%** ($10 / 15$).

Online DEMO

- <https://playground.tensorflow.org>

What is Keras?

Keras: an API for specifying & training differentiable programs

Keras API

TensorFlow / CNTK / MXNet / Theano / ...

GPU

CPU

TPU

Keras is the official high-level API of TensorFlow

- tensorflow.keras (tf.keras) module
- Part of core TensorFlow since v1.4
- Full Keras API
- Better optimized for TF
- Better integration with TF-specific features
 - Estimator API
 - Eager execution
 - etc.

tf.keras

TensorFlow

GPU

CPU

TPU

Who makes Keras? Contributors and backers

 633 contributors





What's special about Keras?

- A focus on user experience.
- Large adoption in the industry and research community.
- Multi-backend, multi-platform.
- Easy productization of models.

Keras vs Tensorflow

Parameters	Keras	Tensorflow
Type	High-Level API Wrapper	Low-Level API
Complexity	Easy to use if you Python language	You need to learn the syntax of using some of Tensorflow function
Purpose	Rapid deployment for making model with standard layers	Allows you to make an arbitrary computational graph or model layers
Tools	Uses other API debug tool such as TFDBG	You can use Tensorboard visualization tools
Community	Large active communities	Large active communities and widely shared resources

Case Study: Classification of Iris Dataset

Iris Dataset

Collected by **Ronald Fisher** in 1936

- Available in **UCI machine learning repository** (<https://archive.ics.uci.edu/ml/datasets/Iris>)
- contains measurements for **150 iris flowers** from three different species. **Each class has 50 observations.**



Iris Versicolor



Iris Setosa



Iris Virginica

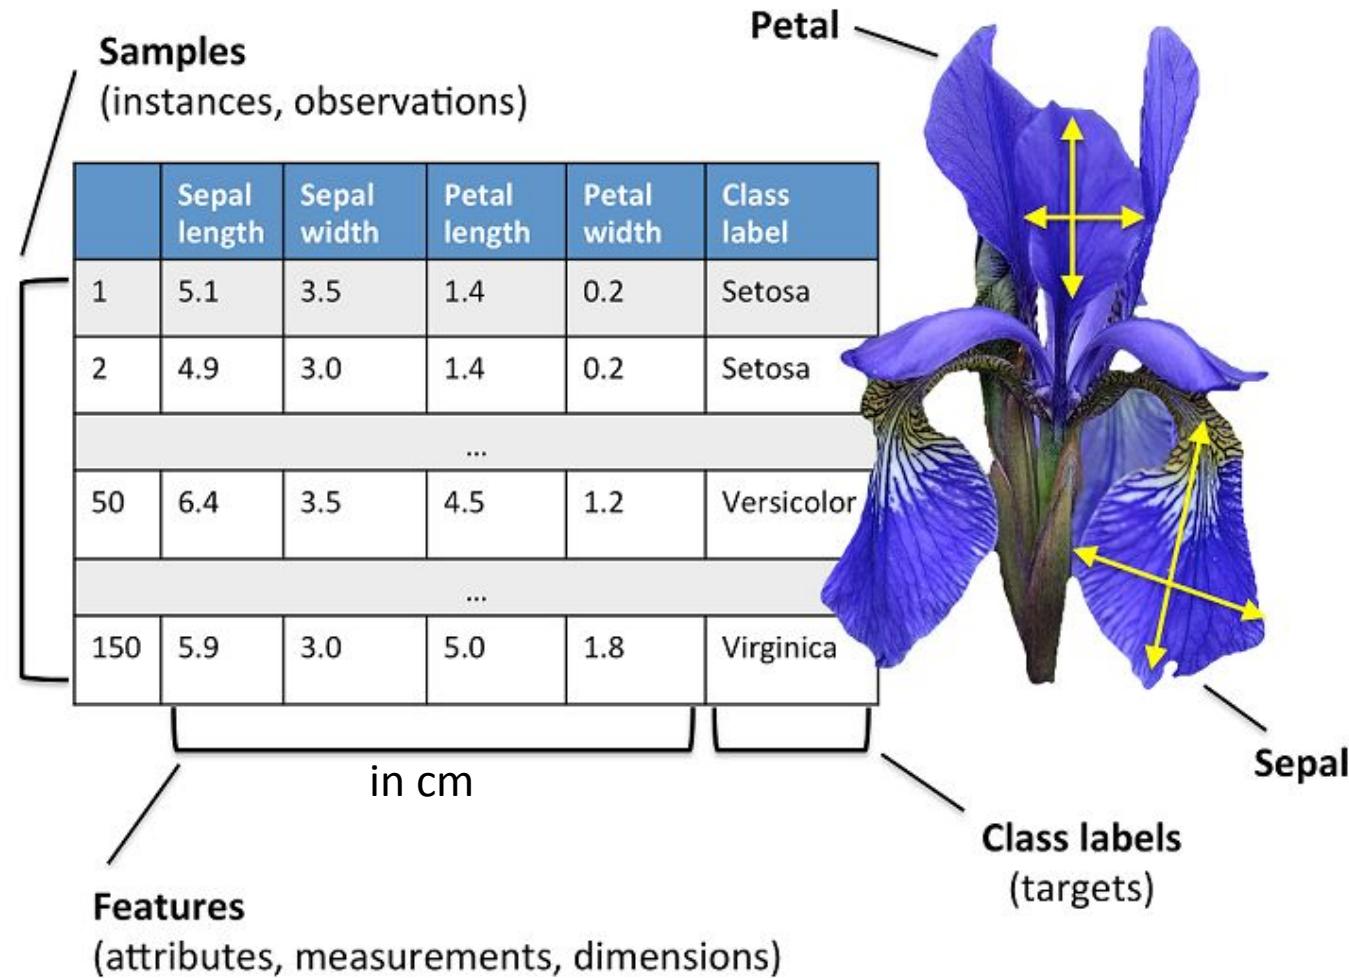
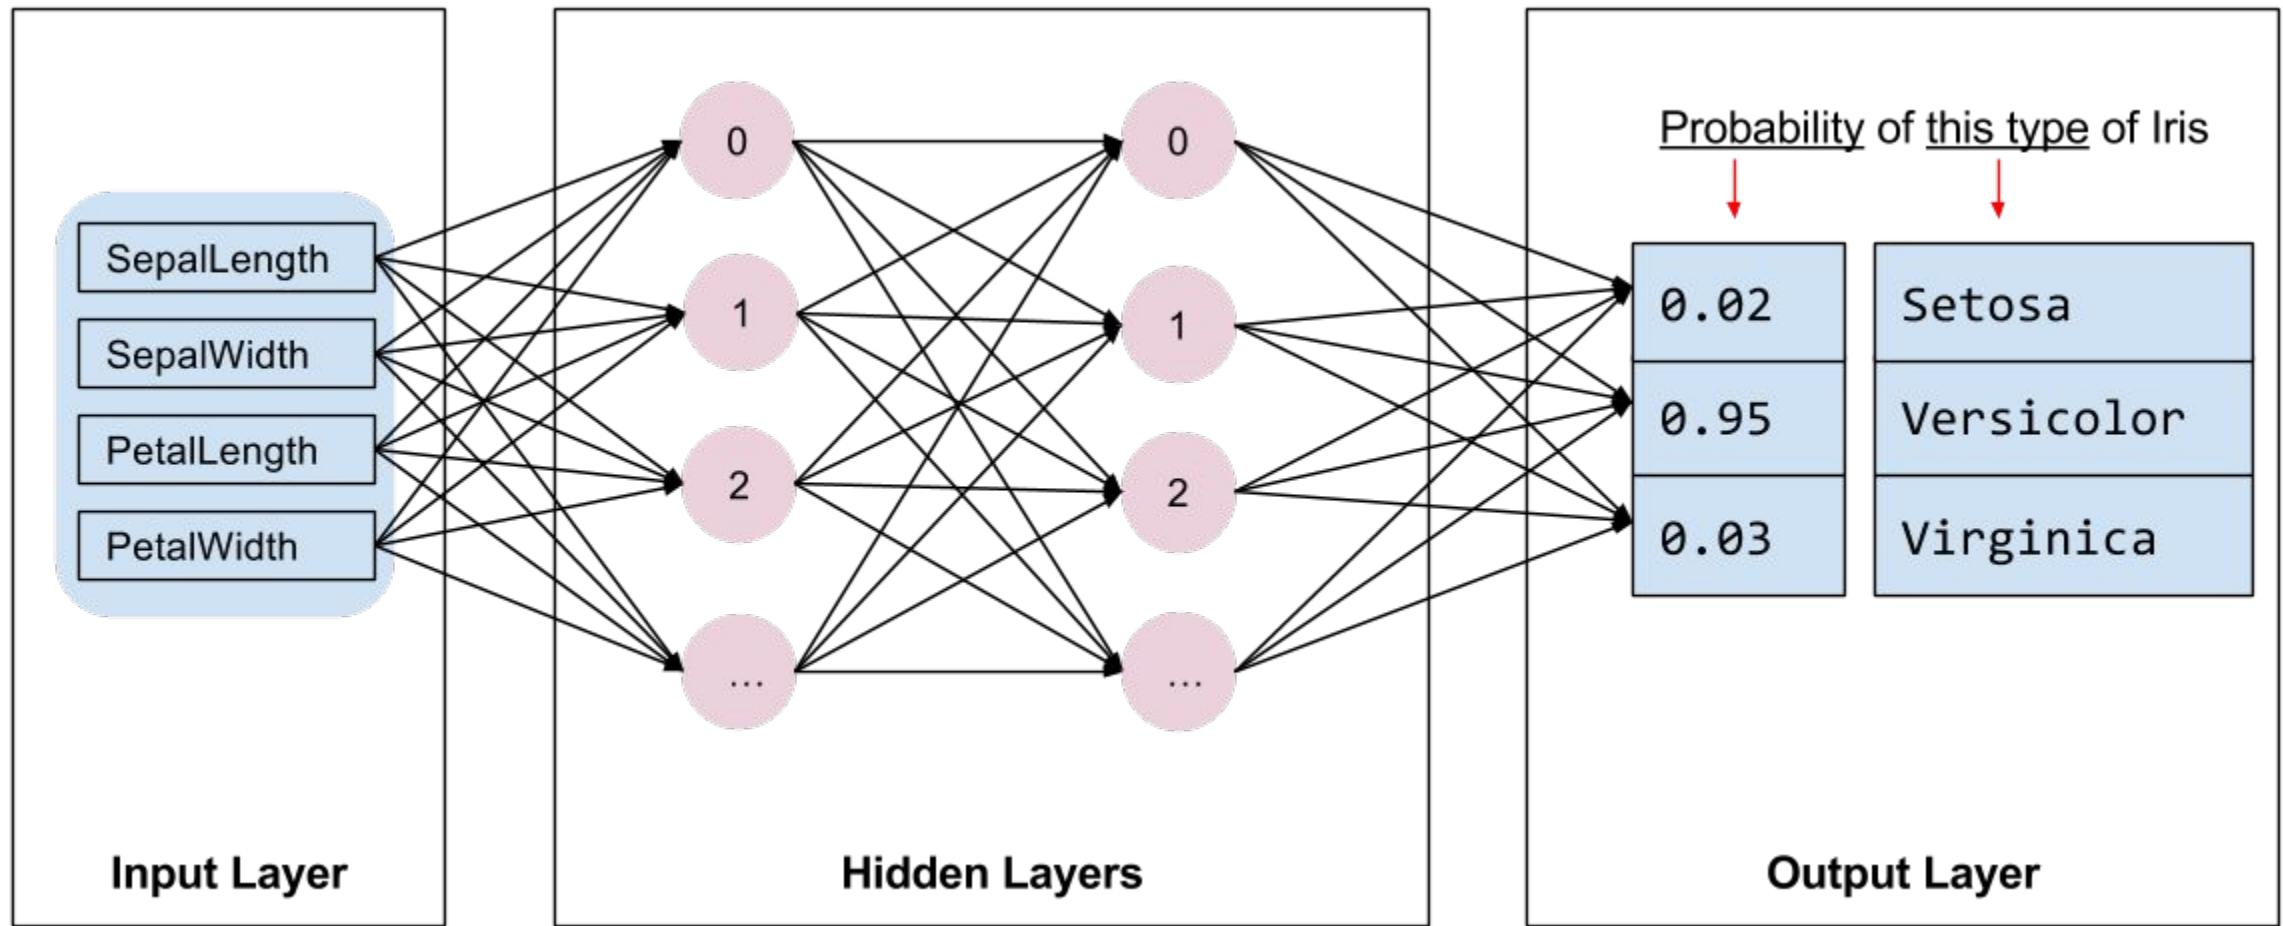


Image Source:

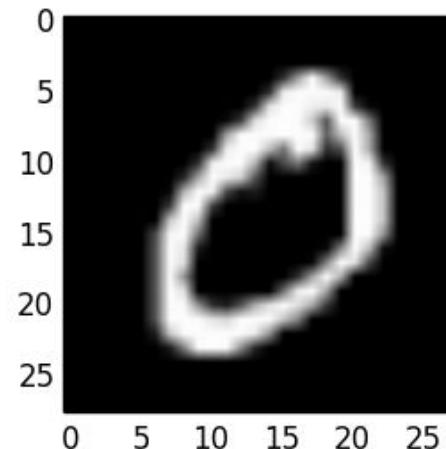
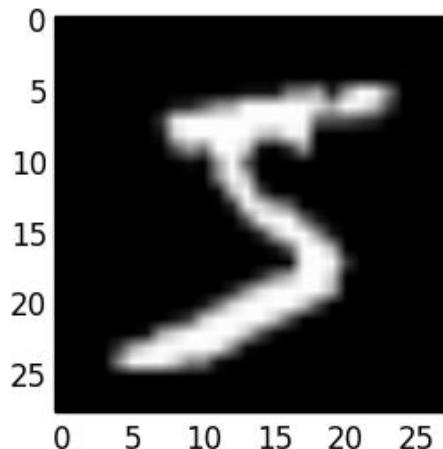
https://selectionpossible.com/Articles/2015_nn_in_3_steps.html

NN Model for Iris classification



MNIST Dataset

- The MNIST database of **handwritten digits**, available from this page, has a **training set of 60,000** examples, and a **test set of 10,000** examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.
- Each Image Size – 28 x 28
- Color Channel – 1 (Black and White)



Input



imread

Blue

Green

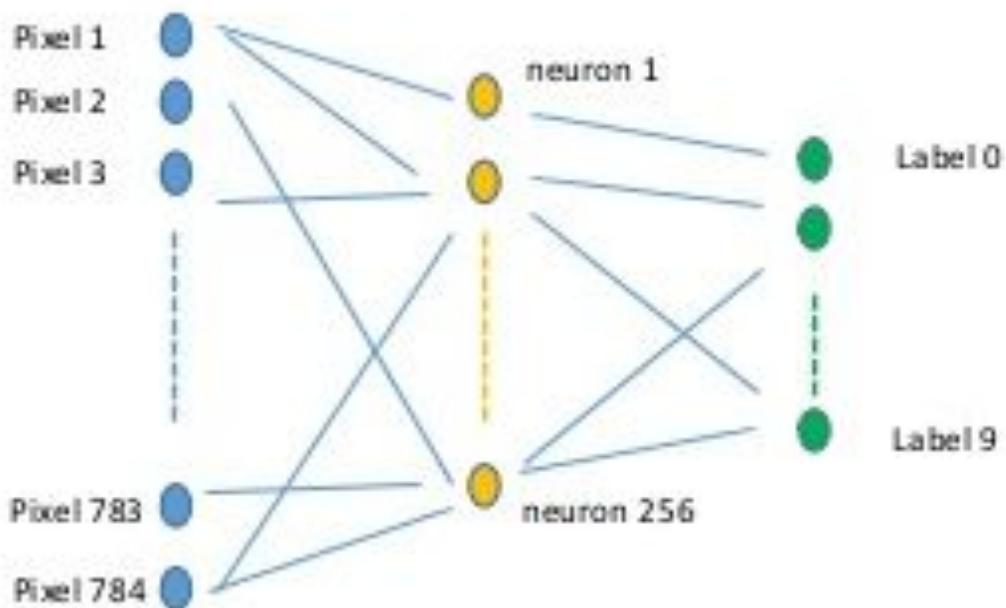
Red

255	134	93	22	2	30	124	142
255	134	202	22	4	30	124	142
123	94	83	2	92	30	124	142
34	44	187	92	34	30	124	142
34	76	232	124	34	30	124	142
67	83	194	202				

reshaped image vector

$$\begin{pmatrix} 255 \\ 231 \\ 42 \\ 22 \\ 123 \\ 94 \\ \vdots \\ 92 \\ 142 \end{pmatrix}$$

Treat image as a vector. It has length 784 (28by28), the number of pixels. One hidden layer (fully connected)



NVIDIA Titan Xp

(Grant by Nvidia Graphics Pvt. Ltd.)

❖ Configuration

- GPU Architecture – Pascal
- NVIDIA Cuda Cores – 3840
- Memory Size – 12 GB GDDR5X
- Memory Speed - 11.4 Gbps
- Nvidia Cuda Compute Capblity – 6.1
- VR Ready



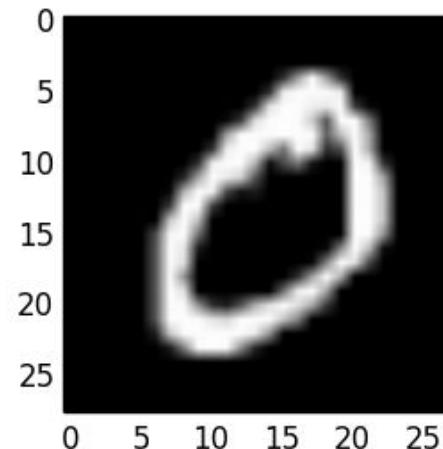
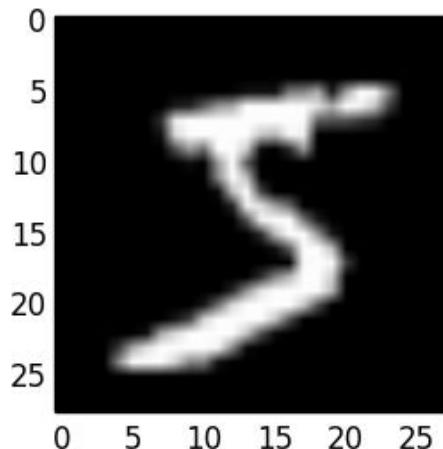
Work Station (Donated by CHARUSAT)

❖ Configuration

- 32 GB RAM
- Core i7 8th Generation CPU
- 2 TB HDD
- ASRock Mother Board Z series
- Cooler Master Cabinet
- 19" FHD LED
- Logitech Keyboard and Mouse

MNIST Dataset

- The MNIST database of **handwritten digits**, available from this page, has a **training set of 60,000** examples, and a **test set of 10,000** examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.
- Each Image Size – 28 x 28
- Color Channel – 1 (Black and White)

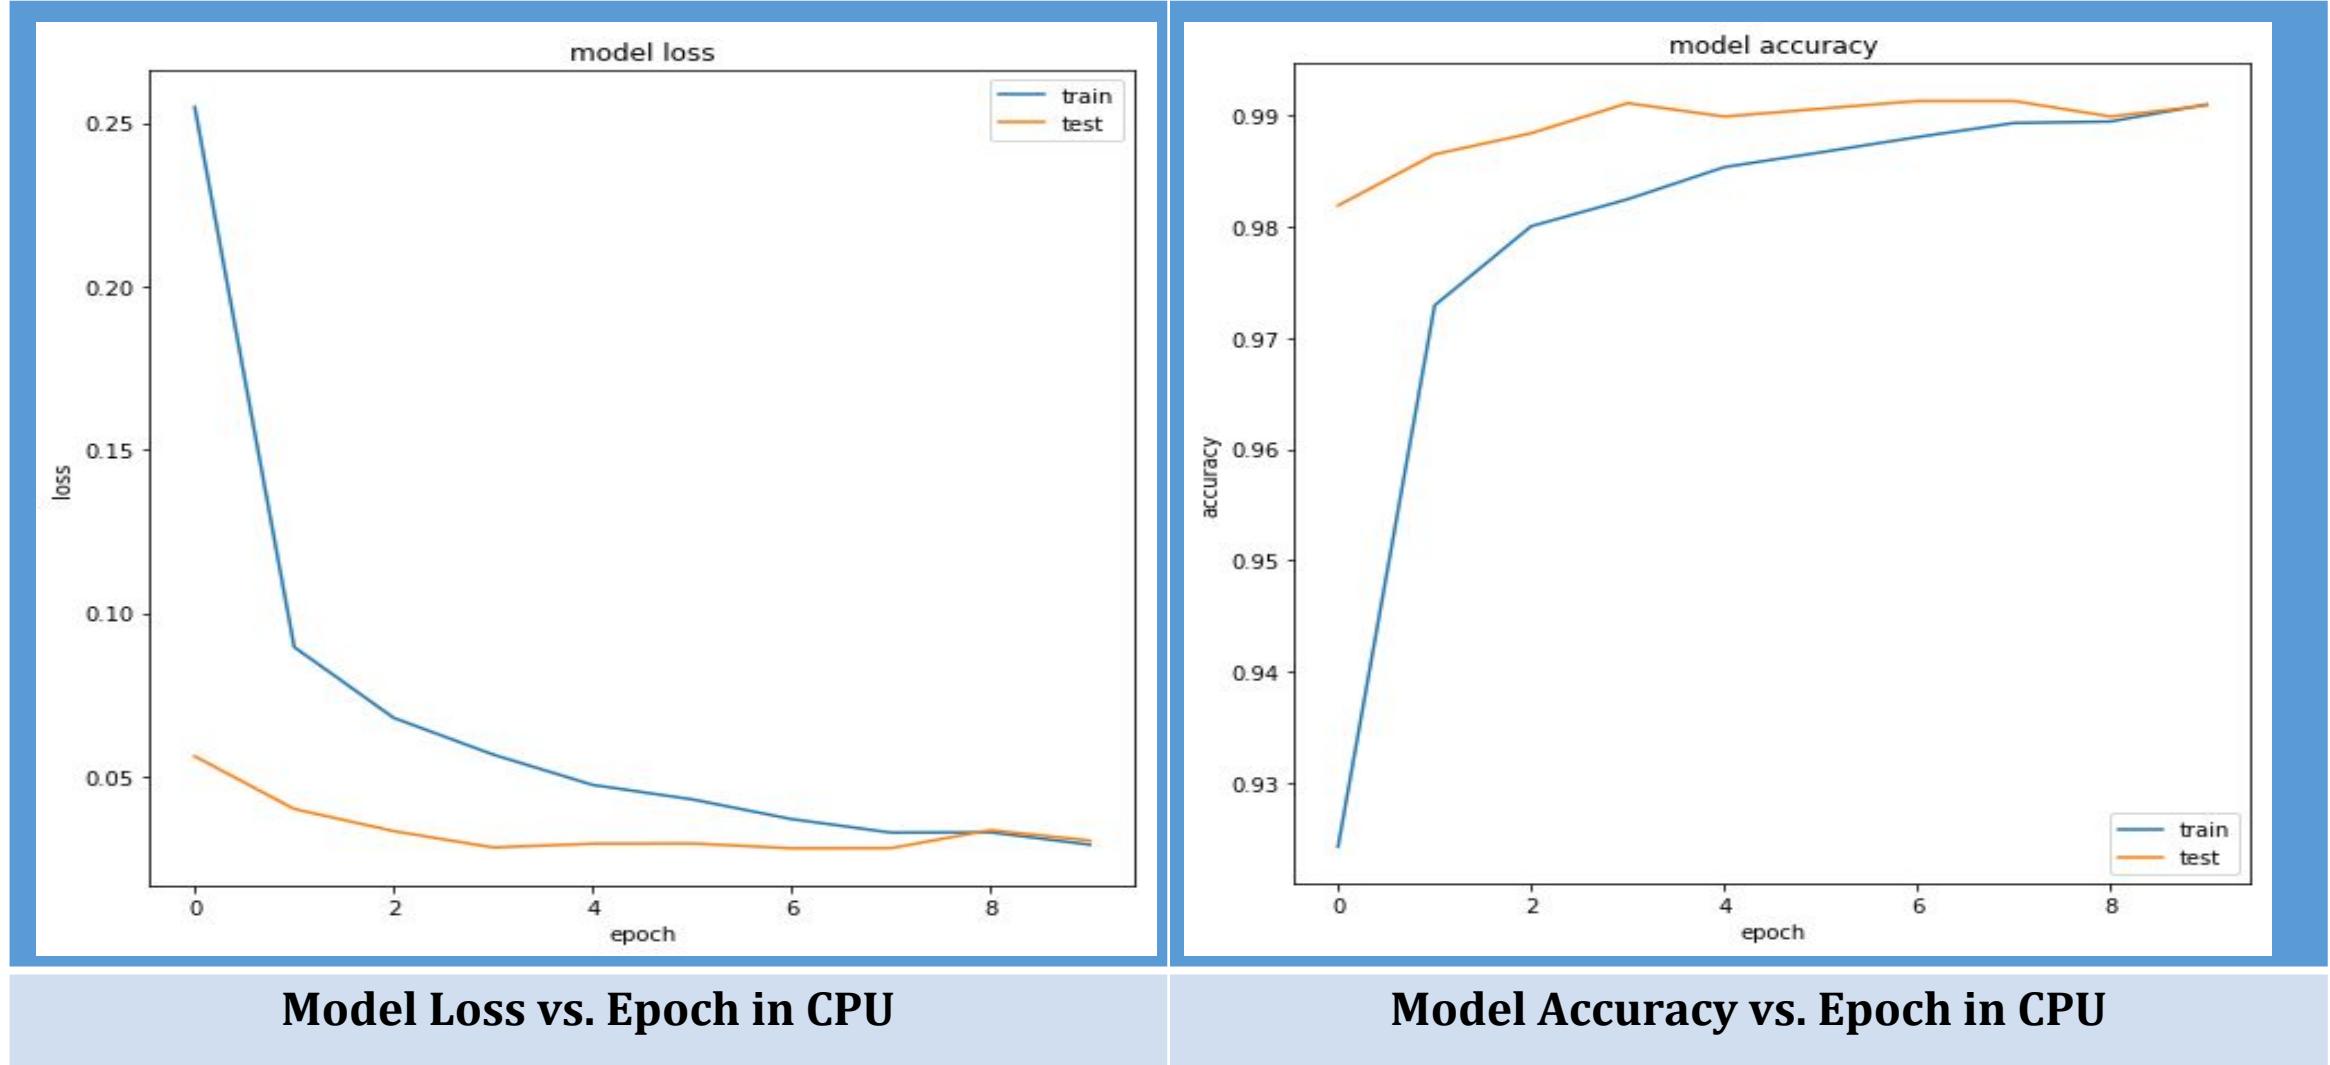


CNN Model

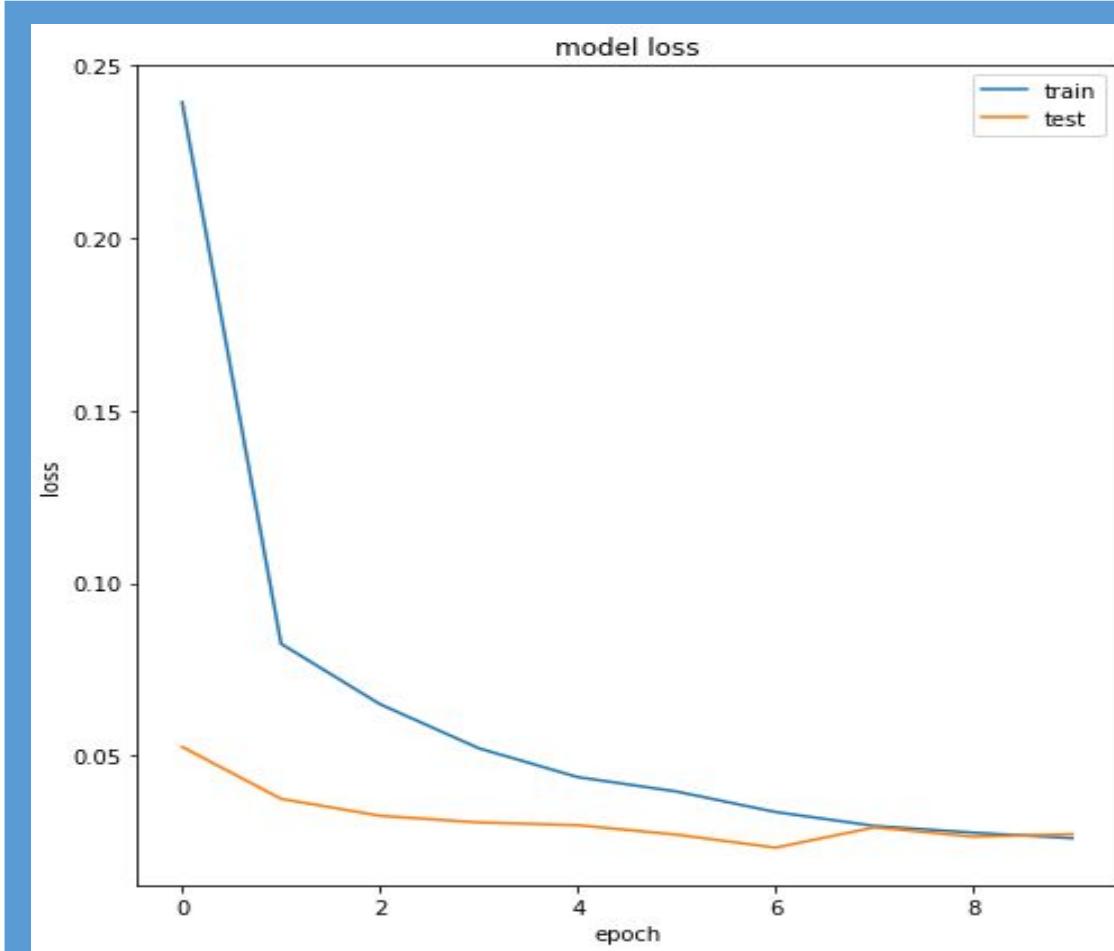
Implementation using Keras
and Tensorflow as backend

Layer (type)	Output Shape	Param #
conv2d_17 (Conv2D)	(None, 26, 26, 32)	320
conv2d_18 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_9 (MaxPooling)	(None, 12, 12, 64)	0
dropout_17 (Dropout)	(None, 12, 12, 64)	0
flatten_9 (Flatten)	(None, 9216)	0
dense_17 (Dense)	(None, 128)	1179776
dropout_18 (Dropout)	(None, 128)	0
dense_18 (Dense)	(None, 10)	1290
=====		
Total params: 1,199,882		
Trainable params: 1,199,882		
Non-trainable params: 0		

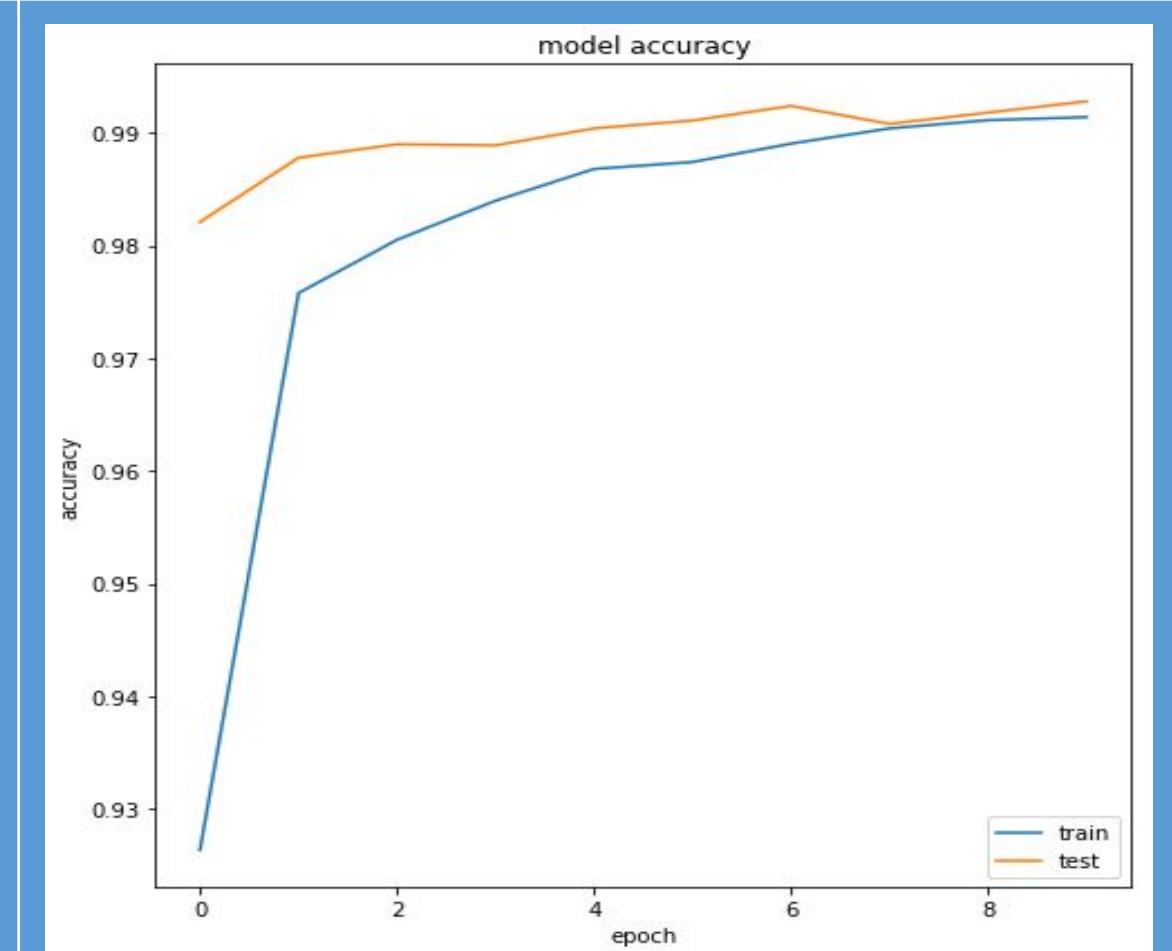
Results in CPU of CNN Model



Results in GPU of CNN Model

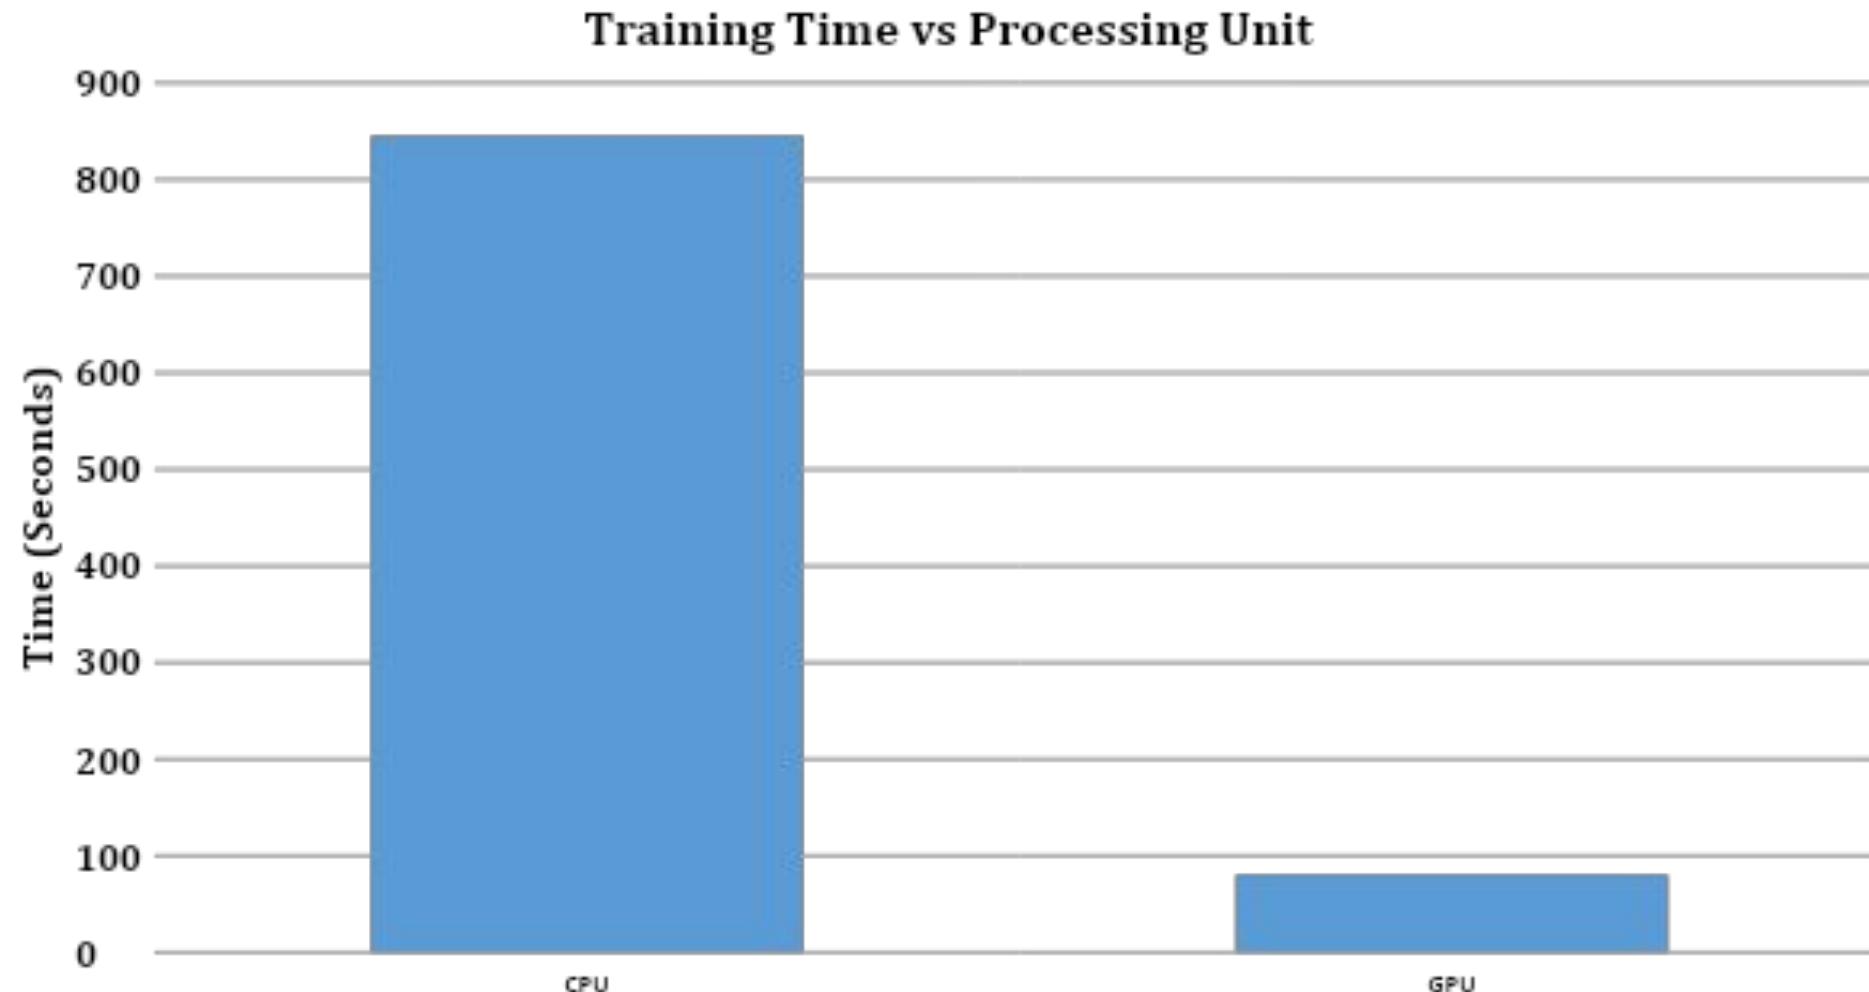


Model Loss vs. Epoch in GPU



Model Accuracy vs. Epoch in GPU

Comparison of training time CNN model in CPU and GPU



ML/DL Conferences/Journals

- **CVPR** - Conference on Computer Vision and Pattern Recognition
- **NIPS** - Neural Information processing systems
- **ICML** - International conference on machine learning
- **ECCV** - European Conference on Computer Vision
- **ICCV** - International Conference on Computer Vision
- **ECML** - European Conference on Machine Learning
- **IJCAI** - International Joint Conference on Artificial Intelligence
- **JML** - Journal of Machine Learning
- IEEE Transactions on Neural Networks
- IEEE Transactions on Pattern Analysis and Machine Intelligence
- Neural Networks
- Artificial Intelligence
- Pattern Recognition
- Expert Systems with Applications
- IEEE Transactions on Knowledge and Data Engineering
- Neural Computing and Applications
- IEEE Transactions on Audio, Speech and Language Processing
- International Journal of Neural Systems

Reading Materials

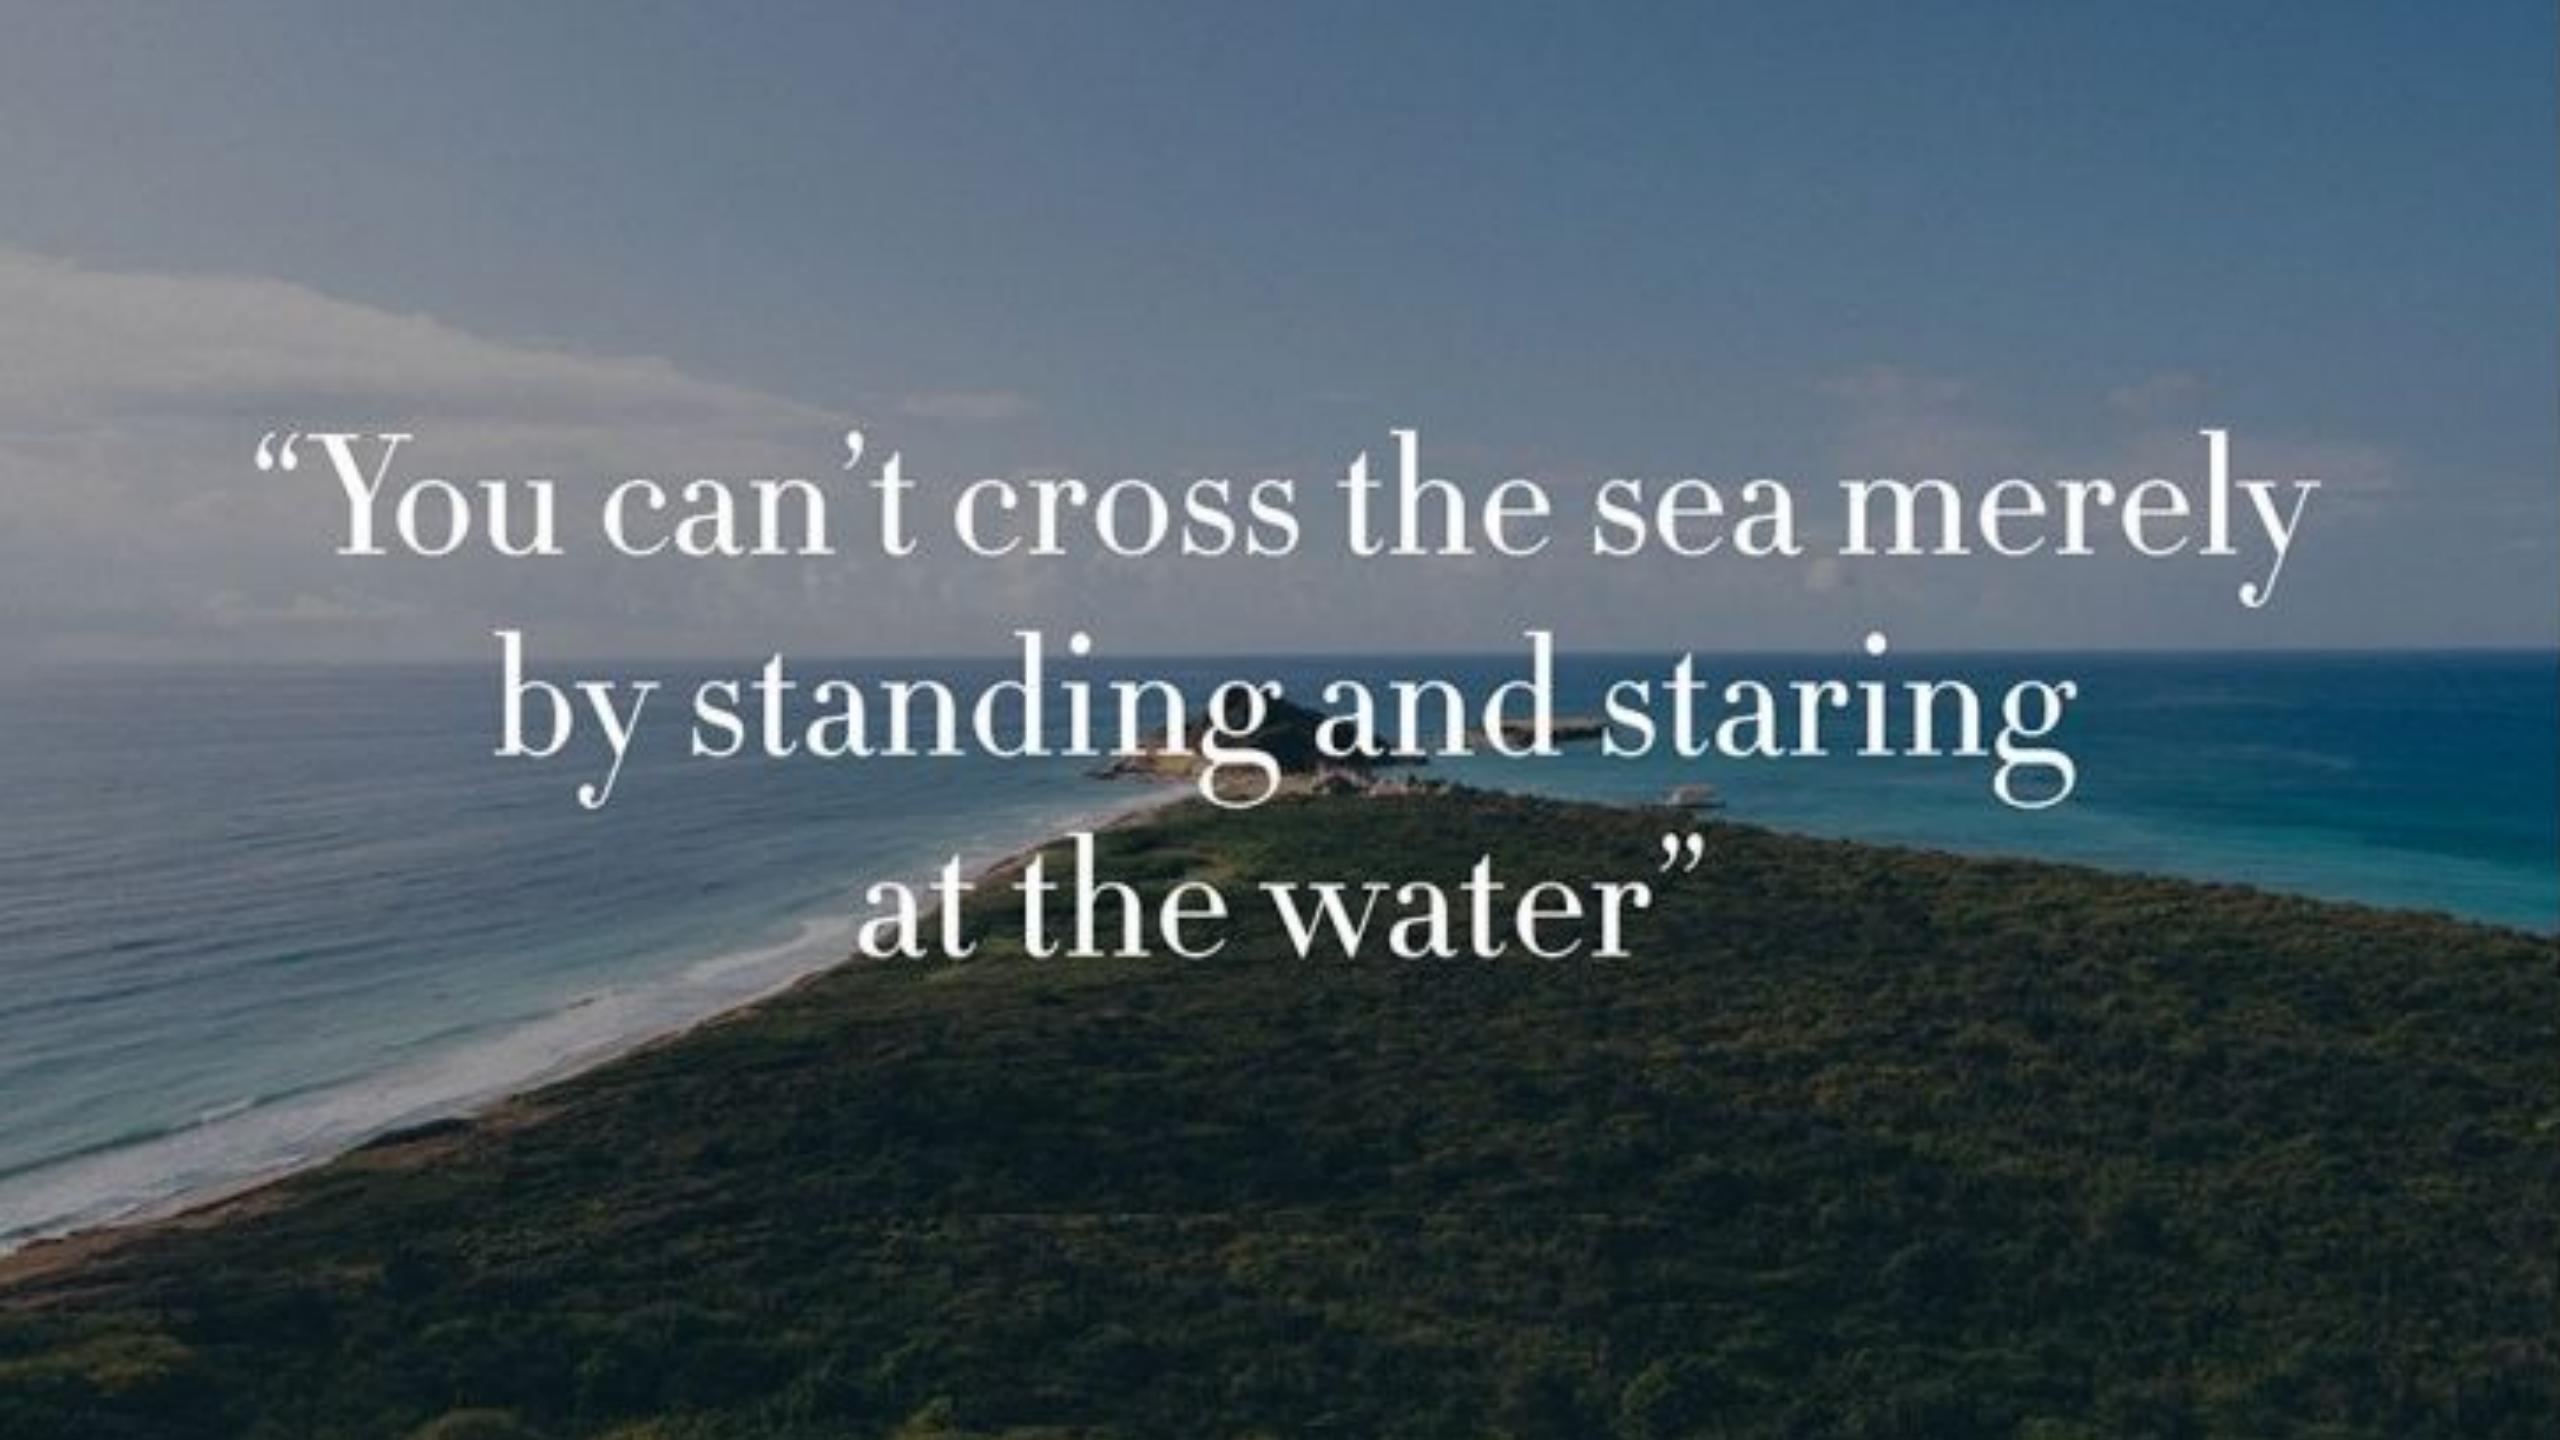
- **Text Book:** <https://www.deeplearningbook.org/>
- <https://www.coursera.org/learn/machine-learning>
- deeplearning.ai
- <http://cs231n.stanford.edu/>
- <https://github.com/floodsung/Deep-Learning-Papers-Reading-Roadmap>
- <https://github.com/terryum/awesome-deep-learning-papers>

Reading Materials

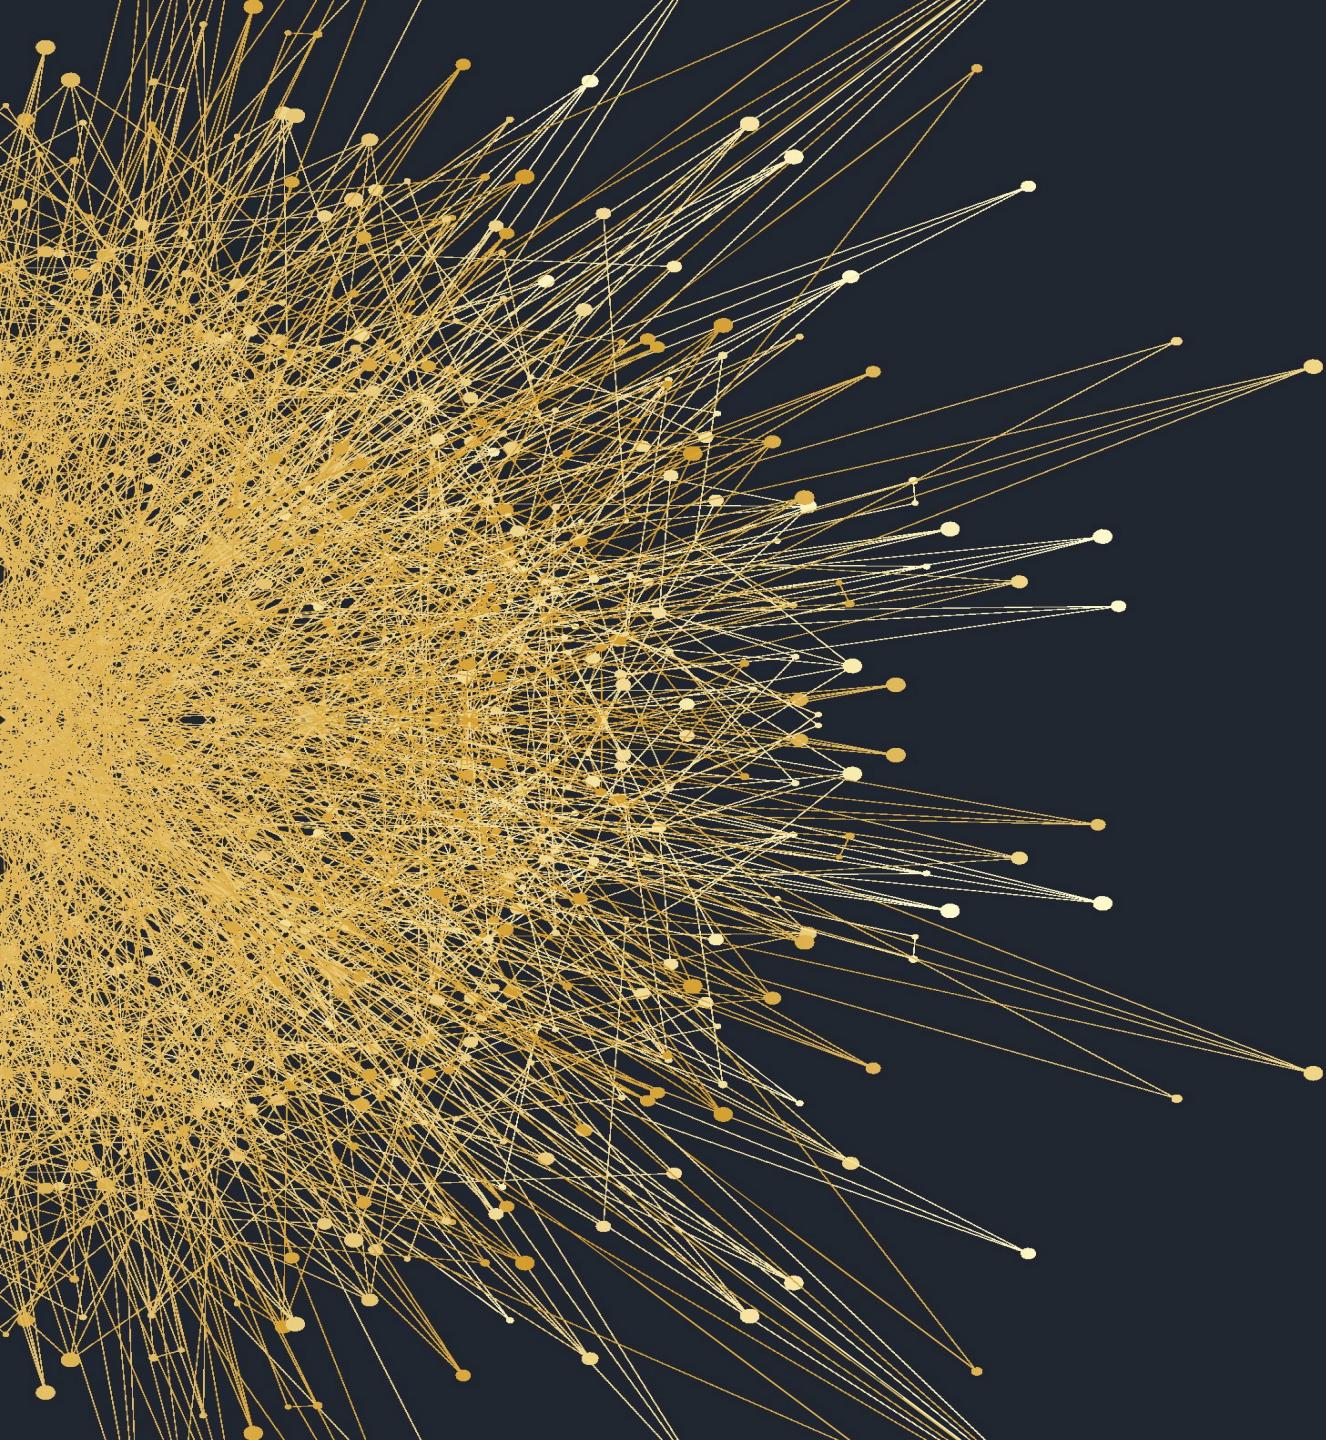
- https://www.tensorflow.org/api_docs/python/
- <https://www.toptal.com/machine-learning/tensorflow-machine-learning-tutorial>
- <https://www.matroid.com/dlwithtf/chap1-2.pdf>
- http://web.donga.ac.kr/yjko/usefulthings/TensorFlow-Basic-Concept_Ko.pdf
- <https://cs224d.stanford.edu/lectures/CS224d-Lecture7.pdf>

References

- Abadi, Martín, et al. “Tensorflow: Large-scale machine learning on heterogeneous distributed systems.” *arXiv preprint arXiv:1603.04467* (2016).
- <https://github.com/tensorflow/tensorflow>
- <https://www.youtube.com/watch?v=yX8KuPZCAMo>
- https://www.youtube.com/watch?v=E8n_k6HNAGs&t=319s

A wide-angle photograph of a coastal scene. In the foreground, there's a steep, dark green hillside covered in dense vegetation. Below the hill, a sandy beach curves along the coastline. The ocean is visible to the right, with its surface appearing calm and slightly blue. The sky above is a clear, pale blue with a few wispy white clouds.

“You can't cross the sea merely
by standing and staring
at the water”



Thank you!

Any

Question ?