PRACTICAL-1

AIM:

To install Hadoop framework, configure it and setup a single node cluster. Use web based tools to monitor your Hadoop setup

IMPLEMENTATION:

The Hadoop framework is written in Java, and its services require a compatible Java Runtime Environment (JRE) and Java Development Kit (JDK).

```
parth642001@parth642001-virtual-machine:~$ sudo apt update
[sudo] password for parth642001:
Hit:1 http://in.archive.ubuntu.com/ubuntu jammy InRelease
Hit:2 http://in.archive.ubuntu.com/ubuntu jammy-updates InRelease
Hit:3 http://in.archive.ubuntu.com/ubuntu jammy-backports InRelease
Get:4 http://security.ubuntu.com/ubuntu jammy-security InRelease [110 kB]
Get:5 http://security.ubuntu.com/ubuntu jammy-security/main amd64 DEP-11 Metada
ta [11.4 kB]
Get:6 http://security.ubuntu.com/ubuntu jammy-security/universe amd64 DEP-11 Me
tadata [608 B]
Fetched 122 kB in 2s (51.3 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
132 packages can be upgraded. Run 'apt list --upgradable' to see them.
```

Apache Hadoop 3.x fully supports Java 8. The OpenJDK 8 package in Ubuntu contains both the runtime environment and development kit.

```
parth642001@parth642001-virtual-machine:~$ sudo apt install openjdk-8-jdk -y
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  ca-certificates-java fonts-dejavu-extra java-common libatk-wrapper-java libat
k-wrapper-java-jni libice-dev libpthread-stubs0-dev libsm-dev libx11-dev
  libxau-dev libxcb1-dev libxdmcp-dev libxt-dev openjdk-8-jdk-headless openjdk-
8-jre openjdk-8-jre-headless x11proto-dev xorg-sgml-doctools xtrans-dev
Suggested packages:
  default-jre libice-doc libsm-doc libx11-doc libxcb-doc libxt-doc openjdk-8-de
mo openjdk-8-source visualvm icedtea-8-plugin fonts-ipafont-gothic
  fonts-ipafont-mincho fonts-wqy-microhei fonts-wqy-zenhei
The following NEW packages will be installed:
  ca-certificates-java fonts-dejavu-extra java-common libatk-wrapper-java libat
k-wrapper-java-jni libice-dev libpthread-stubs0-dev libsm-dev libx11-dev
  libxau-dev libxcb1-dev libxdmcp-dev libxt-dev openjdk-8-jdk openjdk-8-jdk-hea
dless openjdk-8-jre openjdk-8-jre-headless x11proto-dev xorg-sgml-doctools
  xtrans-dev
```

The OpenJDK or Oracle Java version can affect how elements of a Hadoop ecosystem interact. Hence, we need to be specific.

```
parth642001@parth642001-virtual-machine:~$ java -version
openjdk version "1.8.0_312"
OpenJDK Runtime Environment (build 1.8.0_312-8u312-b07-0ubuntu1-b07)
OpenJDK 64-Bit Server VM (build 25.312-b07, mixed mode)
parth642001@parth642001-virtual-machine:~$ javac -version
javac 1.8.0_312
```

Install the OpenSSH server and client.

```
parth642001@parth642001-virtual-machine:~$ sudo apt install openssh-server open
ssh-client -y
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
openssh-client is already the newest version (1:8.9p1-3).
openssh-client set to manually installed.
The following additional packages will be installed:
  ncurses-term openssh-sftp-server ssh-import-id
Suggested packages:
  molly-guard monkeysphere ssh-askpass
The following NEW packages will be installed:
  ncurses-term openssh-server openssh-sftp-server ssh-import-id
0 upgraded, 4 newly installed, 0 to remove and 132 not upgraded.
Need to get 751 kB of archives.
After this operation, 6,046 kB of additional disk space will be used.
Get:1 http://in.archive.ubuntu.com/ubuntu jammy/main amd64 openssh-sftp-server
amd64 1:8.9p1-3 [38.8 kB]
Get:2 http://in.archive.ubuntu.com/ubuntu jammy/main amd64 openssh-server amd64
 1:8.9p1-3 [434 kB]
Get:3 http://in.archive.ubuntu.com/ubuntu jammy/main amd64 ncurses-term all 6.3
-2 [267 kB]
Get:4 http://in.archive.ubuntu.com/ubuntu jammy/main amd64 ssh-import-id all 5.
11-0ubuntu1 [10.1 kB]
Fetched 751 kB in 4s (193 kB/s)
```

Utilize the adduser command to create a new Hadoop user. The username, in this example, is **hdoop**. You are free the use any username and password you see fit. Switch to the newly created user and enter the corresponding password.

```
parth642001@parth642001-virtual-machine:~$ sudo adduser hdoop
[sudo] password for parth642001:
Adding user `hdoop' ...
Adding new group `hdoop' (1002) ...
Adding new user `hdoop' (1002) with group `hdoop' ...
Creating home directory `/home/hdoop' ...
Copying files from `/etc/skel' ...
New password:
Retype new password:
passwd: password updated successfully
Changing the user information for hdoop
Enter the new value, or press ENTER for the default
        Full Name []: Parth
        Room Number []:
        Work Phone []:
        Home Phone []:
        Other []:
Is the information correct? [Y/n] yes
```

The user now needs to be able to SSH to the localhost without being prompted for a password.

```
parth642001@parth642001-virtual-machine:~$ su - hdoop
Password:
hdoop@parth642001-virtual-machine:~$ ssh-keygen -t rsa -P '' -f ~/.ssh/id_rsa
Generating public/private rsa key pair.
Created directory '/home/hdoop/.ssh'.
Your identification has been saved in /home/hdoop/.ssh/id_rsa
Your public key has been saved in /home/hdoop/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:R/s44Pj/pDdsEDMeQu34/dcnkWjlNVH9ctw5/RyWA8w hdoop@parth642001-virtual-ma
chine
The key's randomart image is:
+---[RSA 3072]----+
|          .  o   +|
|         . .  E ..|
|        . o.    ..B|
|        o.*.  oOB|
|        S+o* +.**|
|       o o+o+ + o|
|      . . o+o. ..|
|       .   += o +|
|         ..o+.. o.|
+----[SHA256]-----+
```

Use the cat command to store the public key as **authorized_keys** in the *ssh* directory

Set the permissions for your user with the chmod command

```
hdoop@parth642001-virtual-machine:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized
_keys
hdoop@parth642001-virtual-machine:~$ chmod 0600 ~/.ssh/authorized_keys
hdoop@parth642001-virtual-machine:~$ ssh localhost
The authenticity of host 'localhost (127.0.0.1)' can't be established.
ED25519 key fingerprint is SHA256:KwSx/Dngnw2Cxa7TXe4jhuL8olWAbZTZiOHBJLxHXik.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'localhost' (ED25519) to the list of known hosts.
Welcome to Ubuntu 22.04 LTS (GNU/Linux 5.15.0-41-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

132 updates can be applied immediately.
89 of these updates are standard security updates.
To see these additional updates run: apt list --upgradable


The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
```

Download and extract the Hadoop setup.

```
hdoop@parth642001-virtual-machine:~$ wget https://dlcdn.apache.org/hadoop/commo
n/hadoop-3.3.3/hadoop-3.3.3.tar.gz
--2022-07-21 21:48:48--  https://dlcdn.apache.org/hadoop/common/hadoop-3.3.3/ha
doop-3.3.3.tar.gz
Resolving dlcdn.apache.org (dlcdn.apache.org)... 151.101.2.132, 2a04:4e42::644
Connecting to dlcdn.apache.org (dlcdn.apache.org)|151.101.2.132|:443... connect
ed.
HTTP request sent, awaiting response... 200 OK
Length: 645040598 (615M) [application/x-gzip]
Saving to: ‘hadoop-3.3.3.tar.gz’

hadoop-3.3.3.tar.gz 100%[==================>] 615.16M  1.29MB/s    in 3m 11s

2022-07-21 21:51:59 (3.22 MB/s) - ‘hadoop-3.3.3.tar.gz’ saved [645040598/645040
598]
```

```
hdoop@parth642001-virtual-machine:~$ tar xzf hadoop-3.3.3.tar.gz
hdoop@parth642001-virtual-machine:~$ ls -lrt
total 629936
drwxr-xr-x 10 hdoop hdoop      4096 May  9 23:14 hadoop-3.3.3
-rw-rw-r--  1 hdoop hdoop 645040598 May 11 22:19 hadoop-3.3.3.tar.gz
drwx------  3 hdoop hdoop      4096 Jul 21 21:47 snap
hdoop@parth642001-virtual-machine:~$
```

Hadoop excels when deployed in a fully distributed mode on a large cluster of networked servers. However, if you are new to Hadoop and want to explore basic commands or test applications, you can configure Hadoop on a single node.

This setup, also called pseudo-distributed mode, allows each Hadoop daemon to run as a single Java process. A Hadoop environment is configured by editing a set of configuration files:

- bashrc
- hadoop-env.sh
- core-site.xml
- hdfs-site.xml
- mapred-site-xml

- yarn-site.xml

```
hdoop@parth642001-virtual-machine:~$ sudo nano .bashrc
[sudo] password for hdoop:
```

```
#Hadoop Related Options
export HADOOP_HOME=/home/hdoop/hadoop-3.2.1
export HADOOP_INSTALL=$HADOOP_HOME
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin
export HADOOP_OPTS"-Djava.library.path=$HADOOP_HOME/lib/nativ"
```

```
hdoop@parth642001-virtual-machine:~$ source ~/.bashrc
-bash: export: `HADOOP_OPTS-Djava.library.path=/home/hdoop/hadoop-3.3.3/lib/nat
iv': not a valid identifier
```

```
xid=0 when meet shutdown.
2022-07-21 22:38:19,460 INFO namenode.NameNode: SHUTDOWN_MSG:
/************************************************************
SHUTDOWN_MSG: Shutting down NameNode at parth642001-virtual-machine/127.0.1.1
************************************************************/
```

```
hdoop@parth642001-virtual-machine:~$ ~/hadoop-3.3.3/
-bash: /home/hdoop/hadoop-3.3.3/: Is a directory
hdoop@parth642001-virtual-machine:~$ cd ~/hadoop-3.3.3/sbin
hdoop@parth642001-virtual-machine:~/hadoop-3.3.3/sbin$ ls
distribute-exclude.sh  mr-jobhistory-daemon.sh  start-dfs.sh          stop-balan
cer.sh    workers.sh
FederationStateStore   refresh-namenodes.sh     start-secure-dns.sh   stop-dfs.c
md        yarn-daemon.sh
hadoop-daemon.sh       start-all.cmd            start-yarn.cmd        stop-dfs.s
h         yarn-daemons.sh
hadoop-daemons.sh      start-all.sh             start-yarn.sh         stop-secur
e-dns.sh
httpfs.sh              start-balancer.sh        stop-all.cmd          stop-yarn.
cmd
kms.sh                 start-dfs.cmd            stop-all.sh           stop-yarn.
sh
hdoop@parth642001-virtual-machine:~/hadoop-3.3.3/sbin$ ./start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [parth642001-virtual-machine]
```

Navigate to the *hadoop-3.2.1/sbin* directory and execute the following commands to start the NameNode and DataNode.

Once the namenode, datanodes, and secondary namenode are up and running, start the YARN resource and nodemanagers

```
hdoop@parth642001-virtual-machine:~/hadoop-3.3.3/sbin$ ./start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [parth642001-virtual-machine]
parth642001-virtual-machine: Warning: Permanently added 'parth642001-virtual-ma
chine' (ED25519) to the list of known hosts.
2022-07-21 22:42:11,510 WARN util.NativeCodeLoader: Unable to load native-hadoo
p library for your platform... using builtin-java classes where applicable
hdoop@parth642001-virtual-machine:~/hadoop-3.3.3/sbin$ ./start-yarn.sh
Starting resourcemanager
Starting nodemanagers
hdoop@parth642001-virtual-machine:~/hadoop-3.3.3/sbin$ jps
13920 NameNode
14720 ResourceManager
14833 NodeManager
15177 Jps
14154 DataNode
14459 SecondaryNameNode
hdoop@parth642001-virtual-machine:~/hadoop-3.3.3/sbin$
```

One can also use localhost to access the Hadoop overview.

http://localhost:9870

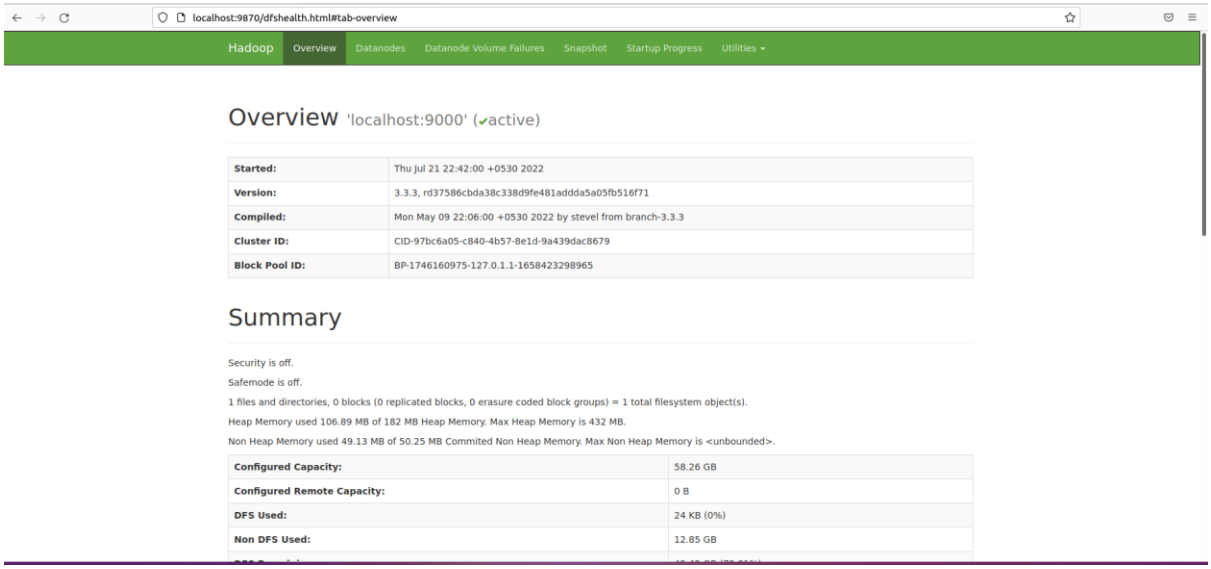

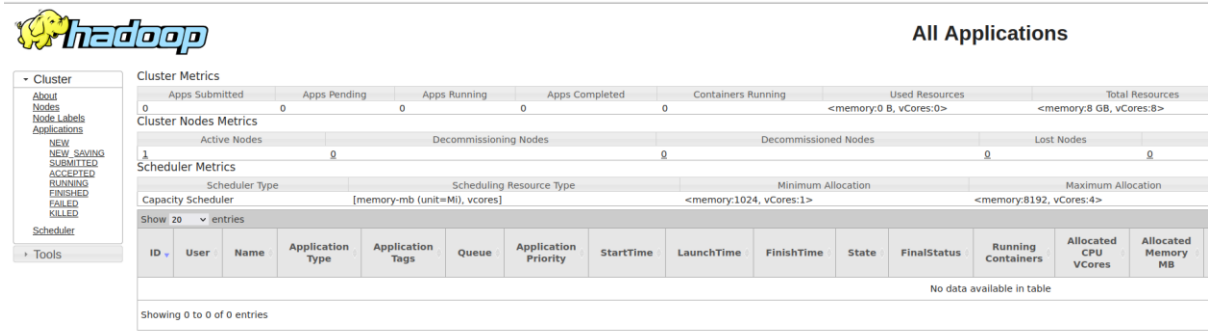

CONCLUSION:

By performing this practical, I learnt how to install and configure Hadoop.