

## PRACTICAL-4

### AIM:

To implement a word count application using the MapReduce API.

### IMPLEMENTATION:

- Firstly, check whether Hadoop is installed or not.

```
hadoop@parth-virtual-machine:~$ hadoop version
Hadoop 3.2.4
Source code repository Unknown -r 7e5d9983b388e372fe640f21f048f2f2ae6e9eba
Compiled by ubuntu on 2022-07-12T11:58Z
Compiled with protoc 2.5.0
From source with checksum ee031c16fe785bbb35252c749418712
This command was run using /home/hadoop/hadoop-3.2.4/share/hadoop/common/hadoop-common-3.2.4.jar
```

- Then, make sure that java compiler is running correctly.

```
hadoop@parth-virtual-machine:~$ javac -version
javac 1.8.0_342
```

- Now, create a folder and a text file for the input.
- Also, create another folder to store java classes files.
- Now, set Hadoop classpath environment variable.

```
hadoop@parth-virtual-machine:~$ export HADOOP_CLASSPATH=$(hadoop classpath)
hadoop@parth-virtual-machine:~$ echo $HADOOP_CLASSPATH
/home/hadoop/hadoop-3.2.4/etc/hadoop:/home/hadoop/hadoop-3.2.4/share/hadoop/common/lib/*:/home/hadoop/hadoop-3.2.4/share/hadoop/common/*:/home/hadoop/hadoop-3.2.4/share/hadoop/hdfs:/home/hadoop/hadoop-3.2.4/share/hadoop/hdfs/lib/*:/home/hadoop/hadoop-3.2.4/share/hadoop/hdfs/*:/home/hadoop/hadoop-3.2.4/share/hadoop/mapreduce/lib/*:/home/hadoop/hadoop-3.2.4/share/hadoop/mapreduce/*:/home/hadoop/hadoop-3.2.4/share/hadoop/yarn:/home/hadoop/hadoop-3.2.4/share/hadoop/yarn/lib/*:/home/hadoop/hadoop-3.2.4/share/hadoop/yarn/*
```

- Create a directory on HDFS.

```
hadoop@parth-virtual-machine:~$ hadoop fs -mkdir /WordCountTutorial
hadoop@parth-virtual-machine:~$
```

- Create another directory in WordCountTutorial for the input.


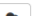

The screenshot shows the HDFS browser interface for the directory `/WordCountTutorial`. It features a search bar, a table of entries, and navigation controls. The table lists one entry named `Input` with permissions `drwxr-xr-x`, owner `hadoop`, group `supergroup`, size `0 B`, and last modified `Sep 01 14:20`. The interface also shows the number of entries (1) and navigation buttons for Previous, 1, and Next.

- Upload the input file to that directory.

```
hadoop@parth-virtual-machine:~$ hadoop fs -put /home/hadoop/Desktop/WordCountTutorial/input_data/input.txt /WordCountTutorial/Input
hadoop@parth-virtual-machine:~$
```

/WordCountTutorial/Input

Go!




Show 

25

 entries

Search:

<input type="checkbox"/>	<div><div></div>Permission</div>	<div><div></div>Owner</div>	<div><div></div>Group</div>	<div><div></div>Size</div>	<div><div></div>Last Modified</div>	<div><div></div>Replication</div>	<div><div></div>Block Size</div>	<div><div></div>Name</div>	
<input type="checkbox"/>	<div>-rw-r--r--</div>	<div>hadoop</div>	<div>supergroup</div>	<div>1.72 KB</div>	<div>Sep 01 14:25</div>	<div>1</div>	<div>128 MB</div>	<div>input.txt</div>	<div></div>

Showing 1 to 1 of 1 entries

Previous

1

Next

- Change the directory to the one where all the files are located.
- Then, compile the java code.

```
hadoop@parth-virtual-machine:~/Desktop/WordCountTutorial$ javac -classpath ${HADOOP_CLASSPATH} -d /home/hadoop/Desktop/WordCountTutorial/classes /home/hadoop/Desktop/WordCountTutorial/WordCount.java
hadoop@parth-virtual-machine:~/Desktop/WordCountTutorial$
```

- Class files are generated in the classes folder.

Home / Desktop / WordCountTutorial / classes			
WordCount.class	WordCount\$IntSumReducer.class	WordCount\$TokenizerMapper.class	

- Put the output files in one jar files.




```
hadoop@parth-virtual-machine:~/Desktop/WordCountTutorial$ jar -cvf wordCount.jar -C classes/ .
added manifest
adding: WordCount$TokenizerMapper.class(in = 1736) (out= 754)(deflated 56%)
adding: WordCount$IntSumReducer.class(in = 1739) (out= 739)(deflated 57%)
adding: WordCount.class(in = 1491) (out= 814)(deflated 45%)
```

- Run the jar file on Hadoop

```
hadoop@parth-virtual-machine:~/Desktop/WordCountTutorial$ hadoop jar /home/hadoop/Desktop/WordCountTutorial/wordCount.jar WordCount /WordCountTutorial/Input /WordCountTutorial/Output
2022-09-01 14:38:17,040 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-09-01 14:38:18,457 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2022-09-01 14:38:18,484 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1662022089534_0001
2022-09-01 14:38:19,146 INFO input.FileInputFormat: Total input files to process : 1
2022-09-01 14:38:19,375 INFO mapreduce.JobSubmitter: number of splits:1
2022-09-01 14:38:20,088 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1662022089534_0001
2022-09-01 14:38:20,090 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-09-01 14:38:21,222 INFO conf.Configuration: resource-types.xml not found
2022-09-01 14:38:21,223 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-09-01 14:38:24,085 INFO impl.YarnClientImpl: Submitted application application_1662022089534_0001
2022-09-01 14:38:24,169 INFO mapreduce.Job: The url to track the job: http://parth-virtual-machine:8088/proxy/application_1662022089534_0001/
2022-09-01 14:38:24,171 INFO mapreduce.Job: Running job: job_1662022089534_0001
2022-09-01 14:38:47,590 INFO mapreduce.Job: Job job_1662022089534_0001 running in uber mode : false
2022-09-01 14:38:47,591 INFO mapreduce.Job: map 0% reduce 0%
```

/WordcountTutorial/Output

Go!














Show 

25

 entries

Search:

<input type="checkbox"/>	 Permission	 Owner	 Group	 Size	 Last Modified	 Replication	 Block Size	 Name	
<input type="checkbox"/>	<a href="#">-rw-r--r--</a>	<a href="#">hadoop</a>	<a href="#">supergroup</a>	0 B	Sep 01 14:39	<a href="#">1</a>	128 MB	<a href="#">_SUCCESS</a>	
<input type="checkbox"/>	<a href="#">-rw-r--r--</a>	<a href="#">hadoop</a>	<a href="#">supergroup</a>	1.58 KB	Sep 01 14:39	<a href="#">1</a>	128 MB	<a href="#">part-r-00000</a>	

- Check the output

```
hadoop@parth-virtual-machine:~/Desktop/WordCountTutorial$ hadoop dfs -cat /WordcountTutorial/Output/*
```

## PROGRAM CODE:

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
```

```
        word.set(itr.nextToken());
        context.write(word, one);
    }
}
}
```

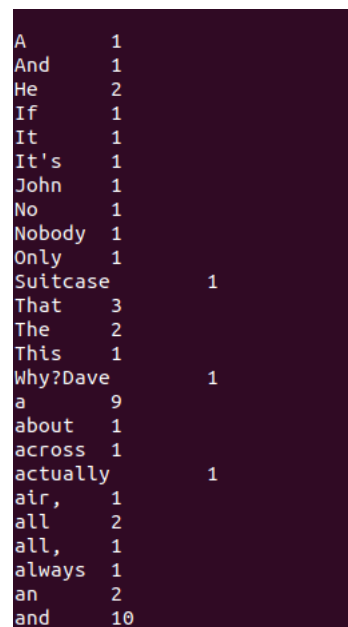
```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
}
```

```
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

## OUTPUT:

A screenshot of a terminal window showing the output of a word count program. The output consists of words and their corresponding counts, separated by a tab character. The words are listed in ascending order of count. The counts range from 1 to 10. The words are: A, And, He, If, It, It's, John, No, Nobody, Only, Suitcase, That, The, This, Why?Dave, a, about, across, actually, air,, all, all,, always, an, and. The counts are: 1, 1, 2, 1, 1, 1, 1, 1, 1, 1, 1, 3, 2, 1, 1, 9, 1, 1, 1, 1, 2, 2, 10.

A	1
And	1
He	2
If	1
It	1
It's	1
John	1
No	1
Nobody	1
Only	1
Suitcase	1
That	3
The	2
This	1
Why?Dave	1
a	9
about	1
across	1
actually	1
air,	1
all	2
all,	1
always	1
an	2
and	10

## CONCLUSION:

In this practical, I learnt to perform wordcount using java, Hadoop and MapReduce Technique.